Certificate-Only Causal Discovery: Monitoring-Valid Edges and Predictive-Fit Orientation

Quang-Vinh Dang[1*]

[1*]British University Vietnam, Hung Yen, Vietnam.

Corresponding author(s). E-mail(s): vinh.dq4@buv.edu.vn;

Abstract

Constraint-based causal discovery deletes an edge when a conditional-independence test fails to reject, so a lack of power is reported as a conclusion, and rerunning it as data arrive accumulates false edges. We study certificate-only discovery, which asserts a feature of the graph only by rejecting a null that is true whenever the feature is absent. The adjacency certificate is the minimum over conditioning sets of sequential tests: a Bayes-factor test of a partial regression coefficient (exact for Gaussian regression errors) or a nonparametric sequential generalised covariance test that allows nonlinear mechanisms (asymptotic). It bounds the probability that a false edge is ever certified, uniformly over time and without faithfulness, within those models. The orientation certificate compares sequentially fitted forecasters for the two factorisations of a pair and reports a direction only when two model classes agree; it certifies better predictive fit, not causation. In 150 random graphs monitored from $\boldsymbol{n = 100}$ to **6400**, CERT-CD reported a false edge in **2.7%** of graphs under Laplace noise (off the Gaussian model), against **27%** to **91%** for PC, GES and DirectLiNGAM and **32%** to **79%** for PC with the same nonparametric test; its recall is low at small samples and most non-adjacent pairs stay undecided. With nonlinear mechanisms the linear test reported a false edge in about half of the graphs and the nonparametric test in **1** to **2** of **60**. On the Tübingen pairs the orientation certificate committed on **12** of **80** scorable pairs and was right on **11**, as often as the twelve most confident pairs of RECI. On **19** untouched recordings of physical devices with known graphs (Causal Chambers) the guarantee did not hold: CERT-CD still reported a false edge, adjusted for dependences the graph cannot explain, in **12** recordings, though with the fewest false edges per recording of the methods compared and a recall of **0.12** to **0.25**. Both certificates fail outside their models, and we report where.

Keywords: causal discovery, anytime-valid inference, e-values, sequential testing, additive noise models, family-wise error control

MSC Classification: 62L12 , 62H22 , 62F03

1 Introduction

A causal discovery algorithm returns a graph, and a practitioner reading it wants to know which of its features the data support. Constraint-based methods such as PC and FCI (Colombo, Maathuis, Kalisch, & Richardson, 2012; Spirtes, Glymour, & Scheines, 2000; Spirtes, Meek, & Richardson, 1995) start from the complete graph and delete an edge whenever a conditional-independence test fails to reject. Failure to reject is compatible with independence, with a dependence too weak to see at the current sample size, and with a test that had little power; the deletion happens regardless. Two consequences follow. Finite-sample correctness requires faithfulness or a quantitative strengthening of it (Robins, Scheines, Spirtes, & Wasserman, 2003; Uhler, Raskutti, Bühlmann, & Yu, 2013), because only such an assumption turns "the test did not reject" into "the variables are independent". And the output is two-valued where the evidence is three-valued: a missing edge means either "the data indicate no edge" or "nothing was established about this pair".

There is a third consequence that matters in practice. Data accumulate, and an analyst who inspects the estimate as they do, and keeps any edge that has ever appeared, is performing optional stopping. A fixed-sample guarantee does not cover this. Section 4.2 measures the effect on eight-node graphs: when PC, GES and DirectLiNGAM are rerun at each sample size from 100 to 6400, between a quarter and nine tenths of the graphs have at least one false edge in the union of the estimates.

This paper studies the opposite discipline, which we call *certificate-only* discovery: assert a feature of the graph only by rejecting a null hypothesis that is true whenever the feature is absent, and never assert anything by failing to reject. An algorithm that obeys the rule starts from the empty graph and adds; it cannot report an edge as absent, only as undecided; and what it controls is the probability of asserting something false. If each test is an anytime-valid process, the guarantee holds at every sample size at once and survives continuous monitoring.

Contributions.

1. *Adjacency certificates under monitoring.* In a structural equation model, d-separation implies that the conditional covariance of a pair given a separating set vanishes, so adjacency has a certifying null that needs no faithfulness. The minimum over conditioning sets of per-set sequential tests bounds the probability that a false edge is *ever* certified (Section 3.4). We give two per-set tests: a Bayes factor for a linear model, exact for Gaussian errors, and a nonparametric sequential generalised covariance test, asymptotic, that allows nonlinear mechanisms (Section 3.4.2). We measure how far from nominal they are off-model (Sections 4.6 and 4.5).

2. *An orientation certificate based on predictive fit.* For a pair of variables we fit forecasters for the factorisations $x \to y$ and $y \to x$ sequentially, and certify a direction when the running mean of the log-score difference has a confidence sequence that excludes zero under two model classes at once (Section 3.5). The claim is about predictive fit and is valid whether or not either model is correct; it identifies the causal direction under additive-noise models.
3. *Evidence under pre-specified protocols.* We compare against PC, GES and DirectLiNGAM under monitoring (Section 4.2), evaluate the orientation certificate on simulated pairs and on the Tübingen cause-effect pairs with a protocol committed before the run (Sections 4.8 and 4.9), evaluate it on physical devices with a known graph (Section 4.10), study scaling with graph size, sample size and the constants of the algorithm (Sections 4.3 and 4.4), and report the regimes where the certificates fail (Sections 4.8 and 4.11).

Scope.

The method targets continuous, observational data from a causally sufficient system whose mechanisms are linear or additive-noise, at sample sizes of a few hundred or more. We chose the benchmarks to match this domain and fixed the selection criteria before running; data outside it are reported as a limitation, not dropped.

Outline.

Section 2 places the work in the literature. Section 3 develops the certificates and the algorithm. Section 4 reports the experiments and Section 5 discusses limitations. Proofs and further results are in the appendices.

2 Related work

Error control in constraint-based discovery.

PC-p (Strobl, Spirtes, & Visweswaran, 2019) bounds an edge's p-value by the maximum p-value over conditioning sets and controls the false discovery rate of edge and orientation decisions with Benjamini–Yekutieli; it is finite-sample but valid at a fixed sample size, and its orientation rules operate within the Markov equivalence class. The conservative PC algorithm of Ramsey, Spirtes, and Zhang (2006) marks ambiguous triples instead of resolving them, and Li and Wang (2009) control the false discovery rate of the edges of PC at a fixed sample size. We split the level over pairs by a union bound; an e-value false-discovery-rate procedure (Wang & Ramdas, 2022) would be less conservative and is not pursued here. Faithfulness-based guarantees break down when parameters are near a cancellation (Robins et al., 2003; Uhler et al., 2013; Zhang & Spirtes, 2008). ϵ-CUT (Shaska & Mitra, 2025) gives one-sided false-edge control without faithfulness at each fixed sample size; Csillag, Struchiner, and Goedert (2025) calibrate batched Fisher-z statistics into per-test e-values inside PC. The certificate-only discipline here is the same one-sided logic made time-uniform and extended to orientation.

Anytime-valid inference.

Ville's inequality (Ville, 1939) underlies confidence sequences (Darling & Robbins, 1967; Howard, Ramdas, McAuliffe, & Sekhon, 2021; Robbins, 1970) and the e-value literature (Grünwald, de Heide, & Koolen, 2024; Ramdas, Grünwald, Vovk, & Shafer, 2023; Vovk & Wang, 2021). Asymptotic confidence sequences (Waudby-Smith, Arbour, Sinha, Kennedy, & Ramdas, 2024) and anytime-valid linear models (Lindon, Ham, Tingley, & Bojinov, 2026) extend the approach to regression coefficients under weak moment conditions, and sequential tests of independence and conditional independence have been built from predictive comparisons and from learned regressions (Podkopaev & Ramdas, 2023; Waudby-Smith, Kennedy, & Ramdas, 2023). Our nonparametric primitive is of the latter kind, built on the generalised covariance measure (Shah & Peters, 2020), whose bias is a product of two regression errors, the structure that debiased machine learning exploits (Chernozhukov et al., 2018) (see also Lundborg, Shah, and Peters 2022 for conditional independence tests in function spaces). What is new here is the use of such tests as certificates, minimised over conditioning sets and combined into a graph. Conditional-independence testing is intrinsically hard: no test with uniform level over all conditional-independence nulls has power against any alternative (Shah & Peters, 2020). Sequential tests therefore either assume a known conditional law (Grünwald, Henzi, & Lardy, 2024; Shaer, Maman, & Romano, 2023) or restrict the null, as we do by testing a partial regression coefficient.

Orientation.

Linear non-Gaussian models identify the full graph (Shimizu, Hoyer, Hyvärinen, & Kerminen, 2006) and DirectLiNGAM (Shimizu et al., 2011) and the pairwise likelihood ratios of Hyvärinen and Smith (2013) estimate it; Shimizu and Kano (2008) test direction using non-normality. Additive-noise models (Hoyer, Janzing, Mooij, Peters, & Schölkopf, 2009; Peters & Bühlmann, 2014) identify direction for nonlinear mechanisms, and location-scale models extend this to heteroscedastic noise (Immer et al., 2023; Xu, Mian, Marx, & Vreeken, 2022). Compression- and quantile-based scores (Marx & Vreeken, 2017; Tagasovska, Vatter, & Chavez-Demoulin, 2020) perform well on the Tübingen pairs (Mooij, Peters, Janzing, Zscheischler, & Schölkopf, 2016). The marginal variances of simulated data can leak the causal order (Reisach, Seiler, & Weichwald, 2021); all our variables are standardised on a warm-up before analysis, which removes that signal. Linear methods fit those pairs poorly (Schkoda & Drton, 2025), and a test that is valid against a misspecified model does not control direction error. Methods that abstain exist: Prakash, Xia, and Erosheva (2024) and Ding and Zhang (2025) return an inconclusive outcome, Prakash, Xia, and Erosheva (2026) gate on a linearity check, and Uehara (2026) attach per-edge resolved or impossible codes in an interactive protocol. None of these gives a time-uniform error statement; ours is a statement about predictive fit that holds without either model being correct (Section 3.5), a comparison in the spirit of Vuong (1989) and of sequential forecaster comparison (Choe & Ramdas, 2024).

3 Method

3.1 Overview

CERT-CD has two certificates and one rule for combining them. The *adjacency certificate* for a pair is a sequential test of the hypothesis that some admissible conditioning set makes the partial regression coefficient zero; it fires when the evidence against that hypothesis, minimised over conditioning sets, is large. The *orientation certificate* for a pair compares two sequentially fitted forecasters of the joint law and fires when one factorisation predicts strictly better. The combining rule is that the graph is built only from fired certificates, starting from the empty graph. Each certified edge carries a bound on the probability that a non-adjacent pair is ever certified (Theorem 7); the bound is exact under Gaussian regression errors and asymptotic otherwise. A certified direction is a statement about predictive fit (Section 3.5); inside the multivariate algorithm the orientation layer has only a conditional guarantee (Appendix B).

3.2 Setting and assumptions

Let $\mathcal{G}^\star$ be a directed acyclic graph on $V = \{1, \dots, d\}$ and let $X = (X_1, \dots, X_d)$ follow a linear structural equation model $X = B^\top X + e$, where B is supported on the edges of $\mathcal{G}^\star$ and the errors $e_1, \dots, e_d$ are mutually independent with finite fourth moments. Observations $o_1, o_2, \dots$ are i.i.d. draws of X and arrive in batches; $\mathcal{F}_t$ is the σ-field of the first t observations. We write $\mathrm{pa}(i)$ and $\mathrm{de}(i)$ for the parents and descendants of i and $\mathcal{S}_k(i,j) = \{S \subseteq V \setminus \{i,j\} : |S| \le k\}$ for the admissible conditioning sets, a family fixed before any data are seen.

Assumption 1 (Separator size) Every non-adjacent pair of $\mathcal{G}^\star$ has a d-separating set of size at most k.

Assumption 1 is structural. It holds whenever k is at least the maximum in-degree of $\mathcal{G}^\star$, because for non-adjacent i, j either $\mathrm{pa}(i)$ or $\mathrm{pa}(j)$ d-separates them; $k = d - 2$ makes it vacuous at an exponential computational cost.

For a pair (i, j) and a set S let $\beta_{i \leftarrow j \cdot S}$ denote the coefficient of X_j in the population regression of X_i on $(1, X_j, X_S)$.

Lemma 1 *If S d-separates i and j in $\mathcal{G}^\star$, then $\beta_{i \leftarrow j \cdot S} = 0$.*

Proof sketch In a linear structural equation model with independent errors the covariance of X is $\Sigma = (I - B)^{-\top} \Omega (I - B)^{-1}$ with Ω diagonal, and the partial covariance of X_i and X_j given X_S is a rational function of the entries of Σ. It depends on the error laws only through Ω. For Gaussian errors, d-separation of i and j by S implies conditional independence by the global Markov property and hence a zero partial covariance (Spirtes et al., 2000); the partial covariance is the same function of (B, Ω) for any other error laws with the same variances, so it vanishes for them as well. The population regression coefficient $\beta_{i \leftarrow j \cdot S}$ is proportional to the partial covariance. □

This is a moment condition, weaker than conditional independence, and it does not require Gaussian errors. Faithfulness is the converse and is not assumed.

3.3 Certificates

Definition 1 Let ϕ be a feature of $\mathcal{G}^\star$. A *certifying null* for ϕ is a hypothesis H_ϕ that holds whenever ϕ is false. A process $(E_t)_{t\geq 0}$ is *Ville-valid* for H if $E_t \geq 0$ and $\mathbb{P}(\exists t : E_t \geq \lambda) \leq 1/\lambda$ for every $\lambda \geq 1$ and every $P \in H$. A *certificate for ϕ at level α* is the event that a Ville-valid process for H_ϕ reaches $1/\alpha$.

Proposition 2 (Soundness) *If ϕ is false, the probability that it is ever certified at level α is at most α.*

The proof is the definition: if ϕ is false then $P \in H_\phi$. The implication runs one way. H_ϕ may also hold when ϕ is true, and the certificate then never fires; this is what buys validity without faithfulness and what costs power. Three consequences follow. There is no certifying null for "i and j are non-adjacent", so the algorithm cannot delete and must start from the empty graph. Its output is three-valued (certified adjacent, certified with a direction, undecided). And power is descriptive: nothing guarantees that a true edge is ever certified.

Ville-valid is weaker than *e-process*, which requires $\mathbb{E}[E_\tau] \leq 1$ at every stopping time of a single filtration. All our bounds use only the event that a process ever reaches a level, so they hold at every data-dependent stopping time of the run, but we do not claim the stronger property. Section 3.4.1 shows why.

3.4 The adjacency certificate

For a pair (i, j) consider

$$H_{ij}^{\text{sep}} : \quad \exists S \in \mathcal{S}_k(i,j) \;\text{ with }\; \beta_{i \leftarrow j \cdot S} = 0. \tag{1}$$

By Lemma 1 and Assumption 1, H_{ij}^{sep} holds for every non-adjacent pair, so it is a certifying null for adjacency. Let $E_t^{(S)}$ be a Ville-valid process for the point null $\beta_{i \leftarrow j \cdot S} = 0$ and define

$$A_t(i,j) = \min_{S \in \mathcal{S}_k(i,j)} E_t^{(S)}(i,j). \tag{2}$$

Proposition 3 *Let $\lambda \geq 1$ and let (i, j) be non-adjacent. Then $\mathbb{P}(\exists t : A_t(i,j) \geq \lambda) \leq 1/\lambda$.*

Proof Choose $S^\star \in \mathcal{S}_k(i,j)$ with $\beta_{i \leftarrow j \cdot S^\star} = 0$. Since $A_t \leq E_t^{(S^\star)}$ pointwise, the event $\{\exists t : A_t \geq \lambda\}$ is contained in $\{\exists t : E_t^{(S^\star)} \geq \lambda\}$, which has probability at most $1/\lambda$. □

That the minimum of e-values is an e-value for a union null is standard (Vovk & Wang, 2021); the proof above needs only the inclusion of events, so the processes $E^{(S)}$

may live in different filtrations. Faithfulness enters only through power: adjacency-faithfulness makes every $E^{(S)}$ grow for a genuinely adjacent pair, and an instance where it fails leaves the pair undecided instead of producing a wrong edge.

3.4.1 The per-set primitive

For S with $|S| = q - 1$ write the regression $X_i = \alpha^\top (1, X_S) + \beta X_j + \varepsilon$ and let S_{ii}, S_{ij}, S_{jj} be the cross-products of the residuals of X_i and X_j after regression on $(1, X_S)$, computed from the running Gram matrix of the centred data, with $p = |S|+1$. With a right-Haar prior on the nuisance parameters and $\beta \mid \sigma \sim N(0, \sigma^2 g)$ for a fixed $g > 0$, the Bayes factor is

$$E_n = (1+G)^{-1/2}\big(1 - \kappa\, \hat{r}_n^2\big)^{-(n-p)/2}, \qquad G = g\, S_{jj}, \quad \kappa = \tfrac{G}{1+G}, \quad \hat{r}_n^2 = \tfrac{S_{ij}^2}{S_{ii} S_{jj}}. \tag{3}$$

It depends on the data only through Schur complements of the Gram matrix, so a single $O(d^2)$ statistic serves every hypothesis (i, j, S). The prior scale g must be fixed before the data: a g that adapts to the running sample size destroys the martingale property. We fix it in advance and use a warm-up prefix, discarded from every process, only to standardise the variables.

Proposition 4 *Suppose $\varepsilon \sim N(0, \sigma^2)$ is independent of (X_j, X_S) and $\beta = 0$. Then (E_n) is a non-negative martingale with mean one in the filtration generated by the sequence of maximal invariants of the nuisance group (translations in the span of $(1, X_S)$ and rescaling of X_i), and therefore $\mathbb{P}(\exists n : E_n \geq \lambda) \leq 1/\lambda$ for every $\lambda \geq 1$.*

The proof is the right-Haar argument of Hendriksen, de Heide, and Grünwald (2021) and Grünwald et al. (2024) and is given in Appendix A. The regressors may have any distribution; Gaussianity is required of the regression error ε, which for $S \supseteq \mathrm{pa}(i)$ with no descendant of i is the structural error e_i.

Remark 1 (The raw-observation filtration) (E_n) is *not* a martingale in the filtration generated by the raw observations. With $S = \emptyset$, $g = 0.25$ and $X_{1:2} = Y_{1:2} = (-1, 1)$, Eq. (3) gives $E_2 = 1$ while $\mathbb{E}[E_3 \mid \mathcal{F}_2] = 1.083 > 1$ under an i.i.d. standard Gaussian null (numerical integration). Amenability of the nuisance group does not repair this. The guarantee used here, a bound on the probability of ever crossing, does not need the raw filtration, which is why Proposition 3 was stated for events.

Remark 2 (Role averaging) Eq. (3) treats X_i as the response, and the roles of i and j are interchangeable under the null. The implementation averages the two role-swapped Bayes factors to remove the arbitrary choice. A crossing of the average implies a crossing of one of the two, each of probability at most $1/\lambda$, so the guaranteed bound for the averaged process is $2/\lambda$.

Remark 3 (Beyond Gaussian errors) When the regression error is not Gaussian, or its variance depends on the regressors, Eq. (3) is the $d = 1$ case of the model-based test of Lindon

et al. (2026), whose guarantee under the moment null $E[W\varepsilon] = 0$ is asymptotic and holds when the model variance equals the heteroscedasticity-robust variance. A robust sandwich version exists (Lindon et al., 2026, Theorem 5.4 of the arXiv version). Finite-sample validity over all finite-variance laws is impossible without tail control (Shah & Peters, 2020). We therefore state the adjacency guarantee as *exact for Gaussian regression errors and asymptotic otherwise*, and Section 4.6 measures the finite-sample departure under heavy tails and heteroscedasticity for both versions.

Remark 4 (One model for both certificates) Both certificates live in the linear structural equation model with independent errors. Adjacency uses only second moments, through Lemma 1; orientation uses the shape of the error laws and, in the additive-noise extension, nonlinear mechanisms, for identifiability only. The validity statement of the orientation certificate (Section 3.5) does not depend on either error law.

3.4.2 A nonparametric primitive

The Bayes factor above needs the regression of X_i on (X_j, X_S) to be linear with a constant-variance error. For the dependences that structural equations with nonlinear mechanisms produce, this is too much to ask: if S contains a mediator between i and j whose effect is nonlinear, a linear regression leaves residual dependence, and the pair is certified although it is non-adjacent (Section 4.5). We therefore give a second primitive, a sequential version of the generalised covariance measure of Shah and Peters (2020), in the spirit of the anytime-valid conditional independence test without the model-X assumption of Waudby-Smith et al. (2023). Write $r_i^S = X_i - \mathbb{E}[X_i \mid X_S]$. If S d-separates i and j then $X_i \perp\!\!\!\perp X_j \mid X_S$ and

$$\mathbb{E}\big[r_i^S \, r_j^S\big] = 0, \tag{4}$$

with no assumption on the form of the conditional means. We use three versions of the product, $r_i^S r_j^S$ and the two products in which one factor is the residual of the square, $r_i^S \, (X_j^2 - \mathbb{E}[X_j^2 \mid X_S])$ and its transpose; each is zero under the separation null, and the squares let the test see dependences that are symmetric, such as $X_j = X_i^2 + \epsilon$, which the plain covariance misses.

Construction.

Let $\hat{f}^S_{v,s}$ be an estimate of $\mathbb{E}[X_v \mid X_S]$ computed from the rows before row s only, for $v \in \{i, j\}$ (and for the squares). The scored product is $z_s = (X_{i,s} - \hat{f}^S_{i,s}(X_{S,s}))(X_{j,s} - \hat{f}^S_{j,s}(X_{S,s}))$. Its running mean has the asymptotic confidence sequence of Waudby-Smith et al. (2024, Theorem 2.2), as in Eq. (6). The set S is *rejected* at the first time at which the sequence for any of the three products excludes zero (level $\alpha_{\text{pair}}/3$ each), and the pair is certified adjacent once *every* $S \in \mathcal{S}_k(i, j)$ has been rejected.

Proposition 5 *Let (i, j) be non-adjacent in $\mathcal{G}^\star$ and let $S^\star \in \mathcal{S}_k(i, j)$ separate them. Suppose the conditional-mean estimates are predictable and the moment conditions of Waudby-Smith et al. (2024) hold for z_s (finite sixth moments of X suffice for the square products), and*

that the product of the estimation errors, $\mathbb{E}[(f - \hat{f}_i)(g - \hat{f}_j)]$*, has a running average that is* $o_P(t^{-1/2})$*. Then the probability that* $S^\star$ *is ever rejected is at most* α_{pair}*, asymptotically, and so is the probability that the pair is ever certified.*

Proof sketch Under (4) the conditional mean of z_s given the past is the product of the two estimation errors, so z_s has a martingale-difference part with mean zero and a predictable remainder; the remainder is negligible by the product-rate condition, and the confidence sequence covers zero uniformly in t with asymptotic probability at least $1 - \alpha_{\text{pair}}/3$ per product. A union bound over the three products gives α_{pair}. The pair is certified only if $S^\star$ is rejected, so its probability is smaller. □

The guarantee is asymptotic, and it is conditional on the estimation error being small, uniformly enough in time that the early fits, which are poor, do not dominate the running average of the bias. No sequential test of conditional independence is valid uniformly over all laws (Shah & Peters, 2020), and we give no regression class and rate for which the product-rate condition holds and do not check it: the additive spline fit can leave a non-vanishing error product if the conditional means are not additive in X_S, in which case a separating set may be rejected and a non-adjacent pair certified. The false-edge rates measured in Section 4 (for example 5 of 60 graphs at $d = 16$ against a level of 0.05) are the evidence on this point, not a proof. In exchange the test covers nonlinear conditional means, which the linear primitive does not.

Implementation.

The regression of each X_v and X_v^2 on X_S is a ridge-penalised additive cubic spline (three knots at the quartiles of the warm-up, inputs clipped at the 1% and 99% warm-up quantiles) with the products of pairs of variables added when $|S| = 2$; it is refit on all rows seen so far when the number of rows has grown by a factor 1.25, and the residuals of the current fit are used for scoring the rows that follow. One multi-output solve per S serves all variables, so a pass costs $O\big(\sum_{|S|\le k} n\, p_S^2\big)$ for the refits and $O(d^2)$ per set per batch for scoring. For $d = 8$ and $k = 2$ a run on 3200 rows takes about 0.3 seconds, and for $d = 20$ and $k = 3$ (about 1350 sets) about 25 seconds on 4000 rows.

3.5 The orientation certificate

Fix a pair (x, y). For each direction we fit a forecaster of the joint law on the past and score it on the next observation. Let $p_s^{x\to y}$ be the density of the forecaster for the factorisation $p(x)\,p(y \mid x)$ fitted on observations $1, \dots, s-1$ and $p_s^{y\to x}$ the mirror forecaster, and set

$$d_s = \log p_s^{x\to y}(x_s, y_s) - \log p_s^{y\to x}(x_s, y_s). \tag{5}$$

Both forecasters are $\mathcal{F}_{s-1}$-measurable, and both are densities on the same space, so d_s is a legitimate score difference whether or not either model is correct. Let $\bar{\mu}_t = t^{-1}\sum_{s\le t}\mathbb{E}[d_s \mid \mathcal{F}_{s-1}]$ and let $\hat{\mu}_t$ and $\hat{\sigma}_t$ be the running mean and standard deviation

of (d_s). The confidence sequence of Waudby-Smith et al. (2024, Theorem 2.2) is

$$\hat{\mu}_t \pm \hat{\sigma}_t \sqrt{\frac{2(t\rho^2+1)}{t^2\rho^2} \log \frac{\sqrt{t\rho^2+1}}{\alpha}}, \tag{6}$$

with ρ tuned to a reference time $t^* = 200$. We certify $x \to y$ at the first t at which the lower end exceeds zero and $y \to x$ when the upper end falls below zero; both are decided only after 60 scored observations.

Proposition 6 *Let $\mathcal{F}_{s-1}$ be the past, let the forecasters be $\mathcal{F}_{s-1}$-measurable, and suppose the martingale differences $d_s - \mathbb{E}[d_s \mid \mathcal{F}_{s-1}]$ satisfy the moment conditions of Waudby-Smith et al. (2024, Theorem 2.2). Then (6) is an asymptotic confidence sequence for $\bar{\mu}_t$ at level $1-\alpha$. In particular, if $y \to x$ predicts at least as well as $x \to y$ on average ($\bar{\mu}_t \leq 0$ for all t), the probability of ever certifying $x \to y$ is at most α asymptotically.*

What is certified is therefore that one factorisation has a strictly better expected log score, and this is valid without either forecaster being correct. It corresponds to the causal direction under the identifiability results for additive-noise models (Hoyer et al., 2009; Peters & Bühlmann, 2014) and for linear non-Gaussian models (Shimizu et al., 2006): if the data come from $x \to y$, the true class is contained in the forecaster family and the estimates are consistent, the expected score of the correct factorisation exceeds that of the reversed one by a positive Kullback–Leibler gap, because for generic mechanisms or non-Gaussian errors the reversed factorisation cannot represent the joint law. Under confounding or selection the better predictor need not be the cause, and the certificate then says nothing about causation. The conclusion of Proposition 6 is asymptotic, and Section 4.8 reports its finite-sample behaviour.

Forecasters.

Both use a generalised-Gaussian density $g(u) \propto \exp(-|u/s|^\kappa)$ with shape $\kappa \in [0.9, 8]$ fitted by maximum likelihood, a generalised-Gaussian marginal for the cause, and a cubic truncated-power spline mean with four quantile knots, set from the warm-up prefix, for the effect. The class `gg` has constant scale; the class `ggh` lets the log-scale be quadratic in the cause, which covers heteroscedastic noise (Immer et al., 2023; Xu et al., 2022). The *agreement* certificate reports a direction only when both classes certify the same one. Each class has the asymptotic guarantee of Proposition 6, so agreement cannot, asymptotically, certify the better-fitting direction wrongly more often than either class alone; this concerns predictive fit and not causation, and the purpose of the rule is protection against model-class mistakes: the constant-scale class alone certifies the wrong direction on every heteroscedastic Gaussian pair it decides in our simulations, and the location-scale class alone is too flexible to resolve linear non-Gaussian pairs (Section 4.8). Preprocessing is fixed in advance: tied values are dequantised, the order is randomly permuted, and $\max(100, \lceil 0.2n \rceil)$ observations form a warm-up prefix used for standardisation and knots.

Algorithm 1 Relative-fit orientation of a pair

Require: $(x_s, y_s)_{s \le n}$, level α, model classes $\{\texttt{gg}, \texttt{ggh}\}$

dequantise ties; permute rows at random; standardise and place knots using the first $m_0 = \max(100, \lceil 0.2n \rceil)$ rows
if $n - m_0 < 100$ **then return** UNSCORABLE
end if
for each class c **do**
  for each batch $s \in \{m_0, m_0 + 50, \dots\}$ **do**
    fit $p^{x \to y}, p^{y \to x}$ in class c on rows $1..s-1$; score the next 50 rows, append the d_s
    if at least 60 scored and no verdict yet **then** compute $\hat{\mu}_t \pm r_t$ from Eq. (6); if the lower end > 0 set verdict $x \to y$, if the upper end < 0 set verdict $y \to x$
    end if
  end for
end for
return the common verdict if both classes agree, else UNDECIDED

Fitting.

Both classes standardise x and y on the warm-up and set four knots at the $20, 40, 60, 80\%$ quantiles of the warm-up. The conditional density of the effect given the cause is fitted by maximising the generalised Gaussian likelihood with the mean a cubic truncated-power spline of the cause, by iteratively reweighted least squares with a continuation over the shape grid, so that the Gaussian shape $\kappa = 2$ and the Laplace shape $\kappa = 1$ are both reachable. In class `ggh` the scale is $\exp(\gamma_0 + \gamma_1 x + \gamma_2 x^2)$, re-estimated by two alternating rounds of weighted refitting. The forecasters are refit on all rows before each batch of 50 scored rows, and each scored row contributes one term d_s. A direction is reported at the first batch after 60 scored rows at which the confidence sequence excludes zero, and the verdict is not revised.

Proof of Proposition 6 The increments d_s are $\mathcal{F}_s$-measurable and the forecasters are $\mathcal{F}_{s-1}$-measurable, so $d_s - \mathbb{E}[d_s \mid \mathcal{F}_{s-1}]$ is a martingale difference sequence. Under the moment conditions, Theorem 2.2 of Waudby-Smith et al. (2024) gives a confidence sequence for the average conditional mean $\bar{\mu}_t$ that holds asymptotically, uniformly in t. If $\bar{\mu}_t \le 0$ for all t then crossing of the lower end above zero implies that the interval fails to cover $\bar{\mu}_t$ at some t. □

What the claim does and does not say.

The proposition is a statement about which forecaster scores better on the sequence of fitted models, not about the data-generating process. If neither class contains the truth, the better-scoring direction can still be the wrong causal one; the certificate is then valid for the statement it makes and misleading as a causal claim. Its causal reading needs identifiability, and then only if the fitted classes approximate the truth well enough for the gap in expected log score to have the right sign. The agreement rule reduces, but does not remove, this risk, and Section 4.8 measures it in simulation.

3.6 The algorithm and its guarantee

CERT-CD maintains a three-valued state for every pair. It streams data in batches, updates the running Gram matrix, recomputes $A_t(i,j)$ for every undecided pair, and certifies adjacency when $A_t(i,j)$ crosses its threshold. The total budget is split as $\alpha = \alpha_A + \alpha_D$, with α_A spread over the $\binom{d}{2}$ adjacency nulls and α_D over the direction nulls, and the thresholds are $\lambda_A = \binom{d}{2}/\alpha_A$ and $\lambda_D = d(d-1)/\alpha_D$, fixed before the data arrive. Which pairs are examined does not enter the bound.

Theorem 7 (Adjacency) *Under Assumption 1, linear structural equations with independent errors, and Gaussian regression errors in (3) for every examined (i,j,S) and both roles of the pair (which holds, for example, when the structural model is jointly Gaussian), the probability that CERT-CD ever certifies an edge between a non-adjacent pair is at most $2\alpha_A$, uniformly over time and at every data-dependent stopping time of the run.*

Proof By Remark 2 each pair's averaged minimum has crossing probability at most $2/\lambda_A$ by Proposition 3. A union bound over at most $\binom{d}{2}$ non-adjacent pairs gives $\binom{d}{2} \cdot 2/\lambda_A = 2\alpha_A$. □

Choosing $\lambda_A = 2\binom{d}{2}/\alpha_A$ gives α_A exactly; the experiments use $\lambda_A = \binom{d}{2}/\alpha_A$, so the guaranteed bound there is $2\alpha_A$. Under non-Gaussian errors the bound holds asymptotically (Remark 3). With the nonparametric primitive of Section 3.4.2 each pair has the level $\alpha_A/\binom{d}{2}$ and no factor 2 arises, so the probability that a false edge is ever certified is at most α_A asymptotically, without linearity or Gaussian errors. The orientation layer inside the multivariate algorithm uses a linear generalised-Gaussian universal-inference certificate whose validity is conditional on the model describing each pair given a block of conditioning variables frozen when the edge is certified (Appendix B); the relative-fit certificate of Section 3.5 is bivariate and is evaluated on pairs. An extension that conditions the relative-fit forecasters on the frozen block is possible but not evaluated here.

3.7 Implementation of the streaming algorithm

Algorithm 2 gives the streaming procedure. The sufficient statistic is the Gram matrix of the standardised rows, updated in $O(bd^2)$ per batch of b rows. Standardisation constants are computed once from the warm-up prefix of m_0 rows, which are then discarded; after that, no quantity used by a monitored process depends on the future.

Complexity.

Each pair has $\sum_{s \le k} \binom{d-2}{s}$ conditioning sets, so one pass over a batch costs $O\big(d^2\, d^k\, k^3\big)$ for the Schur complements, independently of n, and the adjacency state is $O(d^2)$. The direction processes keep the rows since the pair was certified and refit on them, so their cost grows with the time since certification. In our implementation the adjacency certificates alone take about 0.1 second for a run on 3200 rows at $d = 8$, and the full algorithm with orientation about 53 seconds, so the orientation processes dominate the cost (Sections 4.3 and 4.4). Certified pairs are skipped, and the evaluation of A_t stops

Algorithm 2 CERT-CD (streaming)

Require: d, level α, split $\alpha_A = \alpha_D = \alpha/2$, order k, prior scale, batches $B_1, B_2, \ldots$

fix $\lambda_A = \binom{d}{2}/\alpha_A$, $\lambda_D = d(d-1)/\alpha_D$; set all pairs to UNDECIDED
standardise with mean and scale of the warm-up rows; discard them
for each batch B_t **do**
    update the Gram matrix with the standardised rows of B_t
    for each undecided pair (i, j) **do**
        $A \leftarrow \min_{S \in \mathcal{S}_k(i,j)}$ role-averaged $E^{(S)}$ from Eq. (3)
        record $\bar{A}(i, j) \leftarrow \max(\bar{A}(i, j), A)$
        if $\bar{A}(i, j) \geq \lambda_A$ **then** mark (i, j) ADJACENT; freeze $Z_{ij} \leftarrow$ certified neighbours of i and j; open two direction processes
        end if
    end for
    for each open direction process $(a \to b)$ **do**
        extend the numerator on B_t with a forecaster fitted on rows before B_t since opening
        if running maximum of log of the ratio $\geq \log \lambda_D$ **then** mark $a \to b$ certified
        end if
    end for
    output the current marked graph (certified edges, certified arrowheads, all else undecided)
end for

as soon as one conditioning set has evidence below the threshold, since the minimum is then known not to cross.

Choosing the constants.

Validity does not depend on any of the following, and each is fixed before the monitored data arrive. The order k is the only choice with an assumption attached (Assumption 1): any k at least the maximum in-degree suffices, and a smaller k is valid as long as every non-adjacent pair has a separator of size at most k; otherwise the guarantee is lost, as the experiment on violations shows (Section 4.11). The prior scale g sets which effect sizes the test is tuned to detect: with variables standardised on the warm-up, a prior standard deviation of 0.5 ($g = 0.25$) concentrates power on coefficients of a few tenths, and Section 4.4 varies it. The split of α between families affects the thresholds only; the batch size affects how often the process is inspected and nothing else.

Properties of the output.

Because certificates are never retracted, the output is monotone: edges and arrowheads are added and never removed, and so the estimate at time t is contained in the estimate at every later time. Because the guarantee is time-uniform, the run may be stopped at any data-dependent time without changing the stated bounds, and the bound does not depend on the order in which pairs are examined. The guarantee concerns false assertions only. It says nothing about missed edges, and the number of undecided pairs is the price of that choice.

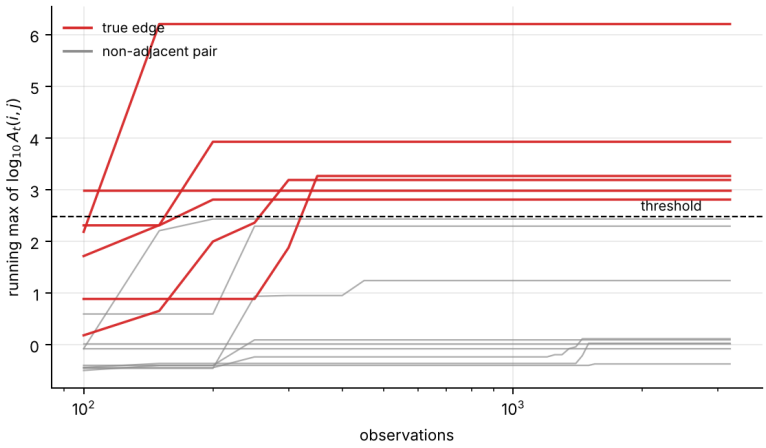


Fig. 1 Running maximum of the adjacency evidence for the 15 pairs of one six-node DAG with Laplace errors; $n = 3200$ observations. Red: true edges. Grey: non-adjacent pairs. Dashed: threshold.

Relation to PC.

CERT-CD examines the same objects as the first layer of PC (partial correlations given small conditioning sets) but reverses the decision rule. PC removes (i, j) as soon as some S yields a non-rejection; CERT-CD adds (i, j) only when *every* S in the family yields rejection, which is why the family size multiplies computation but not the error bound: the bound for a non-adjacent pair uses only the one set that separates it.

A worked example.

Figure 1 follows the adjacency evidence for every pair in one random six-node DAG with Laplace errors (the first draw of a fixed seed, six true edges). The running maximum of $\log_{10} A_t$ is plotted against the number of observations. Evidence for the true edges rises past the threshold of $\lambda_A = \binom{6}{2}/0.05 = 300$ between 100 and 350 observations and is then frozen, since certified pairs are no longer updated; the evidence of the nine non-adjacent pairs stays below it. The highest of these, at about $10^{2.4}$, lies just under the threshold, which is the situation the family bound is designed to cover: with 15 pairs, a few non-adjacent pairs are expected to come close, and the threshold is placed so that, in the Gaussian model, none of them crosses with probability at least $1 - 2\alpha_A$.

4 Experiments

4.1 Protocol

The benchmarks were chosen to match the scope of the method: continuous data, linear or additive-noise mechanisms, a few hundred observations or more. The protocols of the monitoring experiment, the Tübingen evaluation, the simulations of the orientation certificate and of the primitive off the Gaussian model, the held-out pair families and the scaling and nonlinear experiments were written and committed to the repository

Table 1 Probability that a false edge is ever reported over the grid $n = 100, \ldots, 6400$ in 150 random eight-node graphs (Wilson 95% interval for CERT-CD only).

Method	Laplace errors	Gaussian errors
CERT-CD	0.027 $[0.010, 0.067]$	0.040 $[0.018, 0.085]$
PC, $\alpha = 0.05$	0.747	0.773
PC, $\alpha = 0.01$	0.267	0.320
GES-BIC	0.913	0.960
DirectLiNGAM	0.760	1.000

before the runs. The two exceptions are the redesign evaluated on the Causal Chambers (Section 4.10) and the nine design families of Section 4.8, which are flagged where they appear. The baselines IGCI, RECI, ANM-HSIC and the pairwise LiNGAM rule are our re-implementations of the published rules (Appendix D), and PC, GES and DirectLiNGAM are the `causal-learn` implementations with default settings. Data outside the scope are reported in Section 4.11.

4.2 False edges under monitoring

We drew 150 random DAGs on $d = 8$ nodes (edge probability 0.25, coefficients with absolute value in $[0.5, 1.2]$) with Laplace and with Gaussian errors, and ran each method on nested prefixes of $n = 100, 200, 400, \ldots, 6400$ observations. A graph counts as a failure if any prefix's estimate contains an edge that is absent from the true skeleton, which is what an analyst who keeps every edge ever seen would report. CERT-CD ran as one stream with $\alpha = 0.1$ (so $\alpha_A = 0.05$ and a guaranteed bound of 0.1 in the Gaussian model), $k = 2$, warm-up of 50 rows and batches of 100. The comparators were PC (Fisher-z, $\alpha = 0.05$ and 0.01, `causal-learn` defaults), GES with the BIC score, and DirectLiNGAM with edges thresholded at $|B| > 0.05$, each rerun at every n.

CERT-CD reported a false edge in 4 of 150 graphs under Laplace errors and in 6 of 150 under Gaussian errors, against between 27% and 100% for the comparators (Table 1, Figure 2). PC at $\alpha = 0.01$ is close to a Bonferroni correction across the seven looks ($0.05/7 \approx 0.007$) and still reports a false edge in 27% and 32% of the graphs. The comparison is deliberately about monitoring: at a single sample size the comparators' false-edge rates are much lower, for example 42 of 150 graphs for PC at $\alpha = 0.05$ and $n = 6400$ under Laplace errors, so the gap is largely the cost of re-looking, which CERT-CD's guarantee absorbs. The cost is coverage. At $n = 6400$ under Laplace errors CERT-CD's precision is 0.994 and its recall 0.938 ($F_1 = 0.965$), and it leaves 76.8% of the pairs undecided because it never certifies absence (Gaussian: 76.9%). By comparison, PC at $\alpha = 0.01$ has $F_1 = 0.970$ at the final sample size. CERT-CD's orientation is switched on by default here, which splits the budget; edges, not arrowheads, are scored.

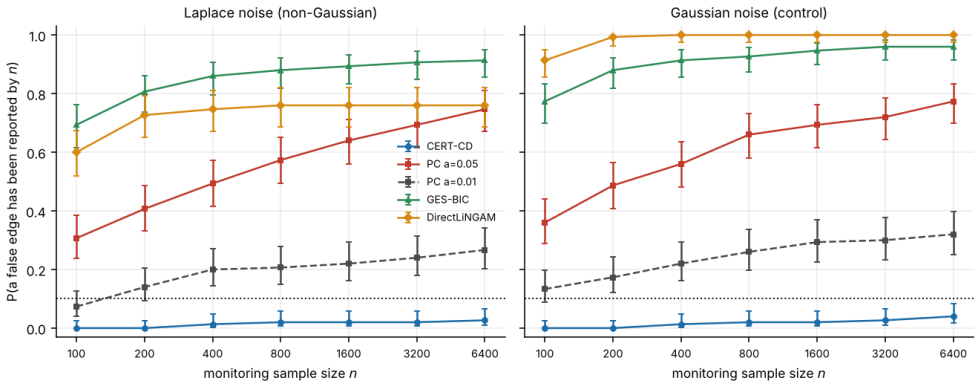


Fig. 2 Cumulative fraction of graphs with a false edge, by sample size, for each method (Laplace and Gaussian errors).

The nonparametric certificate on the same data.

We reran the nonparametric primitive of Section 3.4.2 ($k = 2$, same level) on exactly the same 150 DAGs and prefixes. It reported a false edge at some sample size in 9 of 150 graphs under Laplace errors (0.060, Wilson $[0.032, 0.110]$) and 7 of 150 under Gaussian errors (0.047, $[0.023, 0.093]$), against 4 and 6 for the linear primitive. Both are consistent with the level $\alpha_A = 0.05$; the difference between the two (9 against 4 under Laplace errors, 7 against 6 under Gaussian errors) is within sampling noise. Its skeleton precision and recall at $n = 6400$ are 0.988 and 0.933 under Laplace errors (0.990 and 0.932 under Gaussian), and it has no recall at $n = 100$ (the first rows are used to fit the regressions) and 0.05 at $n = 200$ under Laplace errors, against 0.21 and 0.65 for the linear primitive. For graph sizes $d = 4, 8, 12, 16$ (the scaling data below) it gives 0/60, 3/60, 2/60 and 5/60 graphs with a false edge, with precision ≥ 0.992 and recall between 0.91 and 0.95.

Deleting with the same test.

To separate the effect of the test from the effect of the architecture, we ran PC (depth ≤ 3) with the same nonparametric conditional-independence test used as PC uses a test: at each prefix an edge is deleted when some conditioning set gives a p-value above α (cross-fitted residuals, three products, Bonferroni over the three). On the same 150 DAGs it reported a false edge at some sample size in 115 of 150 graphs at $\alpha = 0.05$ (0.767, Wilson $[0.693, 0.827]$) and 58 of 150 at $\alpha = 0.01$ (0.387, $[0.312, 0.467]$) under Laplace errors, and in 118 and 48 of 150 under Gaussian errors (0.787 and 0.320). At $n = 6400$ under Laplace errors its precision and recall are 0.932 and 0.963 at $\alpha = 0.05$ and 0.974 and 0.954 at $\alpha = 0.01$. The monitoring failure of the comparators is therefore not due to the Fisher-z test or to a particular score: it follows from re-evaluating a delete-on-non-rejection procedure at each sample size.

Table 2 Skeleton precision / recall by sample size, Laplace errors, 150 DAGs at each sample size.

n	100	200	400	800	1600	3200	6400
CERT-CD	1.00/0.21	1.00/0.65	1.00/0.82	1.00/0.87	0.99/0.91	0.99/0.93	0.99/0.94
PC $\alpha = 0.01$	0.99/0.78	0.98/0.85	0.98/0.89	0.99/0.92	0.99/0.94	0.99/0.95	0.99/0.96
PC $\alpha = 0.05$	0.94/0.83	0.95/0.87	0.96/0.90	0.96/0.93	0.96/0.94	0.96/0.95	0.95/0.96
GES-BIC	0.83/0.93	0.84/0.94	0.87/0.96	0.88/0.97	0.92/0.98	0.94/0.98	0.94/0.99
DirectLiNGAM	0.85/0.98	0.91/0.99	0.97/1.00	0.98/1.00	0.99/1.00	0.99/1.00	1.00/1.00

4.3 Sample size and graph size

Power by sample size.

The same runs give precision and recall of the skeleton at each prefix (Table 2, Laplace errors; under Gaussian errors the numbers are similar except for DirectLiNGAM, whose precision is 0.61 to 0.64 at every n). CERT-CD's precision is at least 0.99 at every n, but its recall is low where the sample is small: 0.21 at $n = 100$ and 0.65 at $n = 200$, rising to 0.94 at $n = 6400$. The comparators recover most edges from $n = 100$ and pay with false edges. Under Laplace errors and at a fixed large sample DirectLiNGAM is better than CERT-CD (precision and recall 1.00 at $n = 6400$ against 0.99 and 0.94), because it uses the non-Gaussianity of every variable at once; the claim of CERT-CD is about monitoring, not about accuracy at a single sample size.

Graph size.

We repeated the monitoring experiment for $d = 4, 8, 12, 16$ with 60 DAGs per size, Laplace errors, expected degree about 1.75 (realised mean degree between 1.4 and 1.8) and the grid $n = 100, \ldots, 3200$ (Table 3). The probability that CERT-CD ever reports a false edge is 0, 0.050, 0.067 and 0.083, rising with d but inside the guaranteed bound of 0.1 (which is exact only for Gaussian regression errors); the Wilson upper limit at $d = 16$ is 0.18, so 60 graphs cannot separate this from the bound. For PC at $\alpha = 0.05$, GES and DirectLiNGAM the probability rises to between 0.88 and 1.00 at $d = 12$ to 16, and for PC at $\alpha = 0.01$ to 0.67. CERT-CD's recall stays between 0.92 and 0.96, its precision above 0.99, and the fraction of pairs left undecided grows with d (0.56 at $d = 4$, 0.89 at $d = 16$), because the share of non-adjacent pairs grows and none of them is ever certified absent. The median runtime of one run at $n = 3200$ is 17, 45, 67 and 94 seconds, against under three seconds for each comparator.

4.4 Ablation of CERT-CD

We varied one constant at a time on 20 eight-node DAGs with Laplace errors ($n \leq 3200$, same monitoring criterion; Table 4).

The order k is the only constant that matters for validity. With $k = 0$ and $k = 1$, below the maximal in-degree of many of these graphs, the separator-size assumption is violated and false edges appear in every graph (20/20) and in 15/20 respectively, with precision 0.53 and 0.87; with $k = 3$ nothing changes relative to $k = 2$ except a slightly lower recall. The prior scale and the total level change recall by less than 0.01

Table 3 Monitoring by graph size (60 DAGs per size, Laplace errors, $n \leq 3200$): probability of ever reporting a false edge, with the upper end of CERT-CD's Wilson interval.

Method	$d = 4$	$d = 8$	$d = 12$	$d = 16$
CERT-CD	0.000	0.050	0.067	0.083
Wilson upper limit	0.06	0.14	0.16	0.18
PC, $\alpha = 0.05$	0.233	0.667	0.883	1.000
PC, $\alpha = 0.01$	0.033	0.267	0.367	0.667
GES-BIC	0.317	0.933	1.000	1.000
DirectLiNGAM	0.217	0.867	0.967	1.000
CERT-CD recall at $n = 3200$	0.964	0.930	0.922	0.919

Table 4 Ablation, 20 DAGs, $d = 8$, Laplace errors, monitored to $n = 3200$. Default: $k = 2$, prior scale 0.5, $\alpha = 0.1$, orientation on. Precision and recall are of the skeleton at $n = 3200$.

Variant	ever false edge	precision	recall	median seconds
default	0/20	1.000	0.934	53
$k = 0$	20/20	0.526	1.000	145
$k = 1$	15/20	0.873	0.960	70
$k = 3$	0/20	1.000	0.921	53
no orientation	0/20	1.000	0.934	0.1
prior scale 0.25	0/20	1.000	0.934	52
prior scale 1.0	0/20	1.000	0.927	52
$\alpha = 0.05$	0/20	1.000	0.927	53
$\alpha = 0.2$	1/20	0.993	0.934	56

and leave the number of graphs with a false edge at zero, except that $\alpha = 0.2$ gave one. Switching off orientation leaves the skeleton unchanged, as it should, since the adjacency certificates do not depend on it, and removes nearly all of the computing time: the adjacency certificates alone take about a tenth of a second per run, and the orientation processes dominate the cost.

4.5 Nonlinear mechanisms

The linear primitive assumes linear mechanisms. We generated random DAGs ($d = 8$, edge probability 0.25, 60 per cell, $n \leq 3200$) in which every edge carries a nonlinear function drawn from $\{x, \tanh 2x, \sin 1.5x, x^2 - 1, |x| - 0.8\}$, with Laplace or Gaussian errors, and a linear control. The linear primitive reports a false edge at some sample size in 33 of 60 graphs under Laplace errors and 28 of 60 under Gaussian errors; the nonparametric primitive in 1 of 60 (Wilson $[0.003, 0.089]$) and 2 of 60 ($[0.009, 0.114]$), and it finds more of the true edges (recall 0.83 against 0.66 under Laplace errors, 0.97 against 0.62 under Gaussian errors; precision 0.997 and 0.995 against 0.83 and 0.85).

Table 5 Nonlinear mechanisms, 60 DAGs per row, $d = 8$, $n \leq 3200$: graphs with a false edge at some sample size, and skeleton precision / recall at $n = 3200$.

		linear primitive		nonparametric	
Mechanisms	Errors	ever false	P / R	ever false	P / R
nonlinear	Laplace	33	0.83 / 0.66	1	0.997 / 0.83
nonlinear	Gaussian	28	0.85 / 0.62	2	0.995 / 0.97
linear	Laplace	3	0.99 / 0.93	4	0.99 / 0.92
linear	Gaussian	3	0.99 / 0.93	3	0.99 / 0.93

In the linear control the two primitives agree (false edges in 3 and 4 of 60 graphs under Laplace errors, 3 and 3 under Gaussian errors; recall 0.93 against 0.92; Table 5).

4.6 Calibration of the adjacency primitive

Cost and benefit of the sequential test.

With 2000 replications under the independence null and monitoring up to $n = 1000$, a Fisher-z test inspected at every sample size falsely rejects in 34.6% of runs at nominal $\alpha = 0.05$ (56.5% at 0.10, 81.6% at 0.20; 10.0% at 0.01). The same test at one fixed n rejects 5.6% (0.01: 1.0%), and the e-process 2.4% (0.01: 0.7%). The e-process is conservative by about a factor of two at $\alpha \geq 0.05$ (a factor of 1.4 at 0.01) in this setting, which is the cost of the time-uniform guarantee.

Off the Gaussian model.

The primitive is exact only for Gaussian regression errors. Table 6 reports the realised probability of ever crossing $1/\alpha$ over $n = 50, \ldots, 1000$ (monitored at every observation), with 1000 replications per cell, for a null pair with no covariate effect (S0), a chain $i \leftarrow k \leftarrow j$ tested given k, and a heteroscedastic null, under four error laws. A robust t-statistic confidence sequence (Lindon et al., 2026) is shown for comparison. For Gaussian and Laplace errors without heteroscedasticity the primitive stays below nominal. For skewed or heavy-tailed errors without heteroscedasticity it is up to 2.5 times nominal at $\alpha = 0.01$ (0.025 for lognormal errors), and under heteroscedastic errors it is 2.5 to 3.2 times nominal at $\alpha = 0.01$ and exceeds nominal at $\alpha = 0.05$ for Gaussian (0.069) and Laplace (0.065) errors. The robust version is not uniformly better: it is closer to nominal under heteroscedastic Gaussian and Laplace errors but not under heteroscedastic $t_{2.5}$ and lognormal errors, and it is further from nominal in the heavy-tailed homoscedastic cells (0.098 at $\alpha = 0.05$ for lognormal errors in S0). A separate calibration under six error laws with homoscedastic errors (2000 replications, $n \leq 800$, Appendix D) found every cell below nominal, the largest realised-to-nominal ratio being 0.95 and the median about 0.48. We therefore treat the adjacency bound of the linear primitive as exact in the Gaussian model with homoscedastic errors and as an approximation otherwise, with a departure that is visible under heteroscedasticity and heavy tails.

Table 6 Realised probability of ever crossing, $n = 50, \ldots, 1000$ monitored every observation (primitive, $g = 0.25$) versus a robust-t sequence checked every five observations, 1000 replications per cell.

Scenario	Errors	BF .05	BF .01	robust-t .05	robust-t .01
S0	Gaussian	0.023	0.006	0.044	0.009
S0	Laplace	0.016	0.004	0.032	0.009
S0	$t_{2.5}$	0.035	0.019	0.067	0.021
S0	lognormal	0.043	0.025	0.098	0.035
Chain	Gaussian	0.021	0.008	0.042	0.012
Chain	Laplace	0.018	0.006	0.040	0.011
Chain	$t_{2.5}$	0.026	0.013	0.076	0.035
Chain	lognormal	0.039	0.019	0.086	0.026
Het.	Gaussian	0.069	0.025	0.066	0.022
Het.	Laplace	0.065	0.028	0.053	0.017
Het.	$t_{2.5}$	0.042	0.025	0.077	0.035
Het.	lognormal	0.061	0.032	0.089	0.037

4.7 Graphs that violate strong faithfulness

We generated random linear structural models with Laplace errors ($d = 5$, edge probability 0.35, n up to 3000) until 1000 graphs satisfied δ-strong faithfulness ($\delta = 0.15$) and 500 did not; among the 1780 graphs generated, 56.2% satisfied it. The algorithm ran with $k = 2$ and $\alpha = 0.1$ and its orientation layer on. In the strong-faithfulness stratum CERT-CD certified no false edge (0 of 1000) and made 8 wrong-arrow certifications (rate 0.008); in the other stratum it certified 6 false edges in 500 graphs (0.012, Wilson $[0.006, 0.026]$) and 29 wrong arrows (0.058, $[0.041, 0.082]$; one-sided binomial test against 0.05, $p = 0.23$). PC lost a true edge in 8 of 1000 and in 241 of 500 graphs (0.482, $[0.439, 0.526]$). These are different error types and are not commensurable: a lost edge is an omission and a false edge a commission, and CERT-CD pays for its few commissions with undecided pairs (75% undecided; at $n = 3000$, 97% of the true edges are certified and 94% of the edges oriented, against a CPDAG ceiling of 28% of edges orientable). The orientation layer inside the algorithm is the linear certificate of Appendix B; the count of wrong arrows is a measured rate, not a bound.

4.8 Orientation of pairs

Design families.

The nine families of Table 7 were used while the certificate was being designed: the agreement rule was adopted because each class alone failed on some of them, so the absence of wrong certifications there is not independent evidence. We simulated 50 pairs of $n = 800$ per family and ran the certificate at $\alpha = 0.05$. The agreement certificate committed to a wrong direction in none of the 450 pairs. The constant-scale class certifies the wrong direction on all 32 and 9 pairs it decides in the two heteroscedastic Gaussian families, and the location-scale class is wrong on all 7 it decides for uniform linear errors. Agreement commits on the additive-noise families (41, 18 and 43 of 50

Table 7 Pair orientation, 50 simulated pairs per family, $n = 800$. Entries are (decided, wrong). Forced is the accuracy when a direction is always output.

Family	agree	gg	ggh	forced gg	forced ggh
linear, Laplace	(3,0)	(27,0)	(3,0)	0.96	0.46
linear, uniform	(0,0)	(49,0)	(7,7)	1.00	0.34
linear, Gaussian	(0,0)	(0,0)	(0,0)	0.44	0.44
ANM, tanh	(41,0)	(46,0)	(44,0)	1.00	1.00
ANM, square	(18,0)	(46,0)	(18,0)	1.00	1.00
ANM, sine	(43,0)	(50,0)	(43,0)	1.00	1.00
heteroscedastic Gaussian	(0,0)	(32,32)	(4,0)	0.06	1.00
strongly heteroscedastic	(0,0)	(9,9)	(25,0)	0.64	0.98
heteroscedastic Laplace	(39,0)	(45,0)	(40,0)	1.00	1.00

Table 8 Held-out families, 30 pairs of $n = 800$ each. For the certificates: decided / wrong. Forced-choice columns: accuracy, always outputting a direction.

Family	agree	gg	ggh	IGCI	RECI	ANM	LiNGAM
cubic ANM	6/0	14/0	6/0	1.00	1.00	1.00	0.07
exponential ANM	28/0	30/0	28/0	1.00	1.00	1.00	1.00
logistic ANM	21/0	27/0	21/0	0.37	0.50	1.00	0.00
linear, t_3 errors	0/0	8/0	0/0	0.43	0.37	1.00	1.00
linear, skewed errors	0/0	17/0	12/12	0.10	0.37	1.00	1.00
multiplicative noise	6/0	7/0	17/0	1.00	1.00	0.00	0.53
heteroscedastic, cosine	17/0	26/0	20/0	1.00	1.00	1.00	0.23
post-nonlinear	0/0	4/4	14/0	1.00	0.97	0.00	0.87
bimodal cause, linear	5/5	6/6	10/10	0.40	0.83	1.00	1.00

for the tanh, square and sine mechanisms) and on Laplace heteroscedastic pairs (39 of 50), declines the strongly heteroscedastic Gaussian pairs, and declines linear Gaussian pairs, which are not identifiable.

Held-out families.

To measure behaviour outside the design set we fixed nine further families before running, none used during design, with 30 pairs of $n = 800$ each (Table 8), and compared with four forced-choice published rules (IGCI, RECI, ANM-HSIC and a pairwise LiNGAM rule; our re-implementations, Section 4.1).

The agreement certificate made no wrong certification in eight of the nine families, and wrong certifications on all 5 pairs it decided in the family with a bimodal cause and linear Gaussian noise (5 of 30 pairs, so the overall error rate is 5 of 270 pairs, or 5 of the 83 pairs that it decided). In that family a generalised-Gaussian density cannot represent the mixture marginal of the cause, so the factorisation with the bimodal variable as the effect is fitted better; this is a misspecification of the model class, not a violation of validity, and both classes commit the same error. The agreement certificate

Table 9 Agreement certificate by sample size (decided, of 20 pairs; wrong in parentheses); $n = 800$ has 30 pairs. Also the constant-scale class alone for skewed linear errors.

Family	$n = 200$	400	800 (30)	1600
cubic ANM	0 (0)	0 (0)	6 (0)	15 (0)
heteroscedastic, cosine	3 (0)	6 (0)	17 (0)	19 (0)
linear, skewed errors	0 (0)	0 (0)	0 (0)	0 (0)
constant-scale class alone	0 (0)	1 (0)	17 (0)	18 (0)

commits less often than either class alone (83 of 270 pairs). Among decided pairs it errs in 5 of 83 (6.0%), the constant-scale class in 10 of 139 (7.2%) and the location-scale class in 22 of 128 (17%): the agreement rule removes the errors that the classes do not share (the constant-scale class on post-nonlinear pairs, the location-scale class on skewed linear errors) and not those that they do. Of the forced-choice rules, ANM-HSIC is perfect on seven families and completely wrong on the two with multiplicative and post-nonlinear structure; the others are perfect on some families and at or below chance on others, with no rule dominating. The certificate does not rule out failures in model classes it cannot represent.

Sample size.

For three of the held-out families we varied n (Table 9; 20 pairs per cell, and the $n = 800$ column from Table 8 with 30). The agreement certificate commits more often as n grows and made no wrong certification in any of these cells. The cubic and cosine families go from no or few commitments at $n = 200$ to 15 and 19 of 20 at $n = 1600$. For skewed linear errors it never commits: the constant-scale class alone is right on 18 of 20 at $n = 1600$, but the location-scale class certifies the wrong direction on 15 of 20, so the agreement rule stays silent, which is the protection it is meant to give and also its cost. The four forced-choice rules have their own regimes in these families: ANM-HSIC is right throughout, and the LiNGAM rule fails on the cubic and cosine families.

4.9 The Tübingen cause-effect pairs

The protocol was fixed before the run: all 102 scalar-scalar pairs, official weights, tied values dequantised, a random permutation, at most 2000 rows per pair, $\alpha = 0.05$, one fixed seed. Pairs with fewer than 100 post-warm-up rows (22 of 102, all with at most 194 rows) cannot be scored and count as abstentions; 80 pairs are scored. An earlier linear generalised-Gaussian certificate evaluated on the same pairs decided 14 and was right on 6; the relative-fit certificate was designed afterwards on simulated data, and the repository history shows its constants fixed before this run, but the benchmark was known to us during design.

The agreement certificate decided 12 of 80 scored pairs (15%, which is 11.8% of all 102) and was right on 11 (Table 10); weighted by the official weights its accuracy is 0.816. Forced to always output a direction, our two classes are right on 72.5%

Table 10 Tübingen pairs (80 scored of 102). Accuracy is among decided pairs; interval is Clopper–Pearson 95%; p is the one-sided binomial test of accuracy $> 1/2$.

Certificate	decided	correct	accuracy	95% interval	p
agree (primary)	12	11	0.917	[0.615, 0.998]	0.003
gg only	23	19	0.826	[0.612, 0.950]	0.001
ggh only	18	14	0.778	[0.524, 0.936]	0.015

Table 11 Tübingen pairs: forced-choice accuracy (unweighted / weighted by the official weights) on the 80 scored pairs and, for the baselines, on all 102.

Rule	80 scored pairs	all 102 pairs
CERT-CD constant scale (forced)	0.725 / 0.65	–
CERT-CD location-scale (forced)	0.700 / 0.705	–
IGCI	0.537 / 0.689	0.559 / 0.664
RECI	0.662 / 0.758	0.608 / 0.709
ANM-HSIC	0.650 / 0.618	0.569 / 0.585
pairwise LiNGAM	0.562 / 0.565	0.608 / 0.620
CERT-CD agreement, decided only	11 of 12 correct (0.917)	–

(constant scale, 0.65 weighted) and 70% (location scale, 0.705 weighted) of the 80 scored pairs. Table 11 compares them with the four forced-choice rules computed under the same preprocessing. Among the four rules, RECI is the best on the 80 scored pairs, with 66.2% unweighted accuracy, below the 72.5% and 70% of our classes, and 0.758 weighted, above their 0.65 and 0.705. So the forced predictions of our forecasters are not uniformly better than a simple regression-error rule; the contribution of the certificate is the abstention, which turns a 70–72% forced accuracy on this set into 11 of 12 on the pairs it commits to. Published methods report forced accuracies of roughly 63–83% on all pairs (Immer et al., 2023; Mooij et al., 2016; Tagasovska et al., 2020) under their own protocols; these are not comparable to ours, which are on a subset. We do not compare with abstaining methods at matched coverage. The evidence is modest: 12 decisions, a wide interval (the binomial test treats pairs as independent, which the shared datasets among the Tübingen pairs make doubtful), and a statement about predictive fit.

Sensitivity to the random draws.

The protocol uses one seed for the row permutation, the subsample of long pairs and the dequantisation of ties. We repeated it with three further seeds, keeping everything else fixed. The agreement certificate decided 9, 11 and 10 pairs and was right on 9, 11 and 10 respectively, against 11 of 12 for the frozen seed; the constant-scale class decided 17 to 24 pairs with 14 to 20 correct. The counts are of the same pairs under different random draws and are therefore not independent evidence, but they show

Table 12 Correct / committed at the coverage of the certificate: the certificate itself, and each rule's most confident pairs at the same number k. Tübingen: 80 scored pairs; held-out simulated families: 270 pairs; chamber edges: 94 pair-runs.

Pairs	Coverage (k)	certificate	RECI	IGCI	ANM-HSIC	LiNGAM
Tübingen	agree (12)	11/12	11/12	4/12	9/12	9/12
	constant scale (23)	19/23	21/23	10/23	16/23	14/23
held-out	agree (83)	78/83	82/83	73/83	67/83	57/83
	constant scale (139)	129/139	130/139	118/139	109/139	90/139
chambers	agree (31)	31/31	31/31	31/31	29/31	30/31
	constant scale (48)	48/48	48/48	48/48	43/48	46/48

that the headline is not an artefact of one lucky permutation. The one error of the frozen seed (pair 47) did not recur.

Comparison at the same coverage.

The certificate abstains, the four rules do not, so we also ranked the pairs by each rule's own confidence score (a margin of the statistic that the rule compares) and took the k most confident, with k equal to the number of pairs on which the certificate commits (Table 12). On the Tübingen pairs RECI is right on 11 of its 12 most confident pairs, as many as the agreement certificate, and on 21 of 23 and 16 of 18 against 19 of 23 and 14 of 18 for the constant-scale and location-scale classes; on the held-out simulated pairs it is right on 82 of 83 against 78 of 83; on the chamber pairs RECI and IGCI are right on all of them, as are the classes. At matched coverage, then, the certificate is not more accurate than RECI ranked by confidence. What the certificate adds is a stopping rule: it decides how many pairs to commit to from the data, with an asymptotic statement about the probability of a wrong better-fit direction, whereas a top-k list needs the coverage specified in advance and carries no statement about its error. The rules are our implementations (Appendix D); LOCI and QCCD (Immer et al., 2023; Tagasovska et al., 2020) are not included, since we did not obtain a trustworthy implementation.

Figure 3 shows two pairs: pair 5, the first on which the agreement certificate commits, correctly, and pair 47, the only pair on which it commits and is wrong; the wrong decision is made on 254 observations, with 154 scored, and the confidence sequence only just excludes zero.

4.10 Real physical data with a known graph: the Causal Chambers

The Causal Chambers (Gamella, Peters, & Bühlmann, 2025) are two physical devices, a wind tunnel and a light tunnel, whose causal graphs are known from the construction of the devices and which provide real measurements. We used two recorded experiments in which the actuators are randomised, chosen before any result was seen: the wind tunnel with random intake loads (10 varying variables, 8 true edges among them) and the light tunnel with uniform inputs (20 variables, 39 true edges). A third experiment, the wind tunnel with random loads on both fans, was dropped before running because the two load variables are identical in the recording, so the two nodes

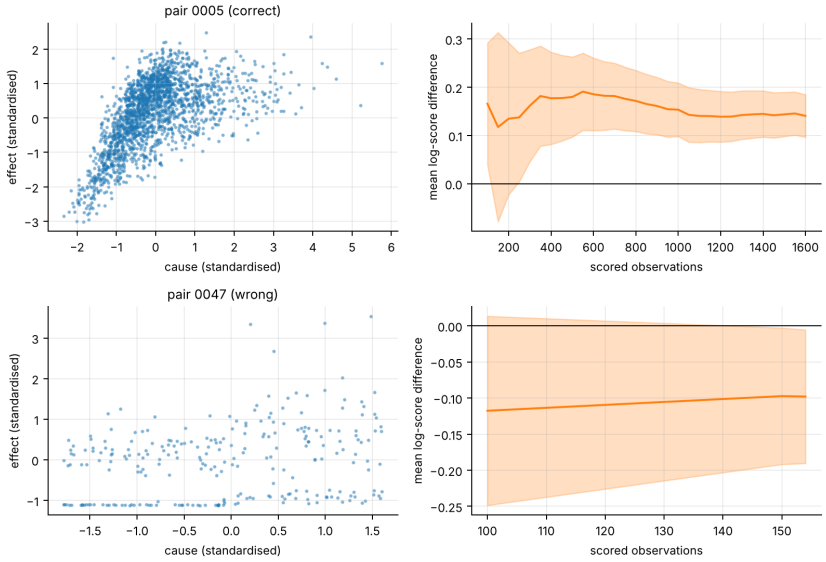


Fig. 3 Tübingen pairs 5 (top) and 47 (bottom): scatter plot and confidence sequence for the running mean log-score difference under the constant-scale class; a lower end above zero certifies the true direction.

cannot be told apart. Variables are the nodes of the chamber graph with at least 20 distinct values; no node is hidden. For each dataset we drew 10 random subsamples of 4000 rows (rows are shuffled, to break the serial dependence of the recording) and ran each method on nested prefixes $n = 250, \ldots, 4000$, as in Section 4.2. The light tunnel recording has 10000 rows, so its ten subsamples overlap heavily and are not independent; counts out of ten there are indicative only.

Linear certificate and comparators.

With the linear certificate at $k = 2$, and with all four comparators, a false edge (an edge of the algorithm's output that is not an edge of the chamber graph) was reported in every one of the 10 subsamples of both devices (rows "strict" of Table 13). CERT-CD's precision at $n = 4000$ was 0.40 and 0.41, similar to PC's, and DirectLiNGAM was the most accurate on the light tunnel (0.79 precision, 0.59 recall). These results motivated the changes below, so what follows is a redesign made after seeing these data, not a test on fresh ones.

Sources of the false edges.

Inspecting the persistent false edges showed three things. First, in the light tunnel the separator of ir_1 from ir_2 is the five parents of ir_1, larger than $k = 2$, so Assumption 1 fails; raising k to 3 removes all but three pairs (the three infrared and visible sensor pairs described below). Second, the mechanisms are nonlinear, which the nonparametric primitive of Section 3.4.2 handles. Third, some pairs the graph calls non-adjacent are dependent given the graph's own separating set. To make this precise without

Table 13 Causal Chambers, 10 subsamples of 4000 rows each. Columns: "ever", subsamples (of 10) with a false edge, adjusted for V, at some sample size; "adj.", false edges per subsample at $n = 4000$, counting a certified pair as correct if it is a graph edge or in the reference set V; "strict", the same counting graph edges only; "rec.", recall of the graph edges (8 and 39). Wind tunnel: $d = 10$, $|V| = 11$; light tunnel: $d = 20$, $|V| = 5$. The CERT-CD variants have orientation off.

	wind tunnel				light tunnel			
Method	ever	adj.	strict	rec.	ever	adj.	strict	rec.
CERT-CD nonparam., $k = 3$	1	0.1	5.9	0.12	0	0.0	3.0	0.18
CERT-CD lin., $k = 3$	10	1.1	3.7	0.39	4	0.4	3.5	0.18
CERT-CD nonparam., $k = 2$	1	0.2	6.2	0.12	10	5.6	9.6	0.19
CERT-CD lin., $k = 2$	10	1.1	4.6	0.39	10	5.9	9.9	0.18
PC, $\alpha = 0.05$	10	2.0	7.9	0.53	10	7.9	11.9	0.24
PC, $\alpha = 0.01$	10	1.8	7.2	0.44	10	6.4	10.4	0.22
GES-BIC	10	8.3	16.8	0.69	10	8.9	12.9	0.60
DirectLiNGAM	10	4.6	10.8	0.61	10	2.9	6.1	0.59

using any of the algorithms under test, we computed, for every non-adjacent pair, the generalised covariance statistic given the parents of the non-descendant (additive spline regression) on up to 20000 rows (the whole recording, 10000 rows, in the light tunnel), and call the pairs with $|z| > 5$ the *reference set* V. It has 11 pairs in the wind tunnel: the six pairs of pressure sensors, which are mutually dependent given their parents with residual correlation 0.999, the two rotation speeds ($z = 11$), and four pairs of the intake rotation speed with a pressure sensor ($|z| \approx 5.6$, residual correlation about -0.03, close to the cutoff). It has 5 in the light tunnel: three pairs of the infrared and visible sensors of the same unit (residual correlation 0.75 to 0.80) and two pairs with $|z| \approx 6$ and residual correlation about -0.07. These are dependences that the graph cannot explain, plausibly a shared source or drift that the graph does not contain. A certified edge that is neither a graph edge nor in V is called a *false edge, adjusted*; the scoring is partly circular for any method that detects conditional dependence with large n, and we report both versions.

Redesigned certificate.

With the nonparametric primitive and $k = 3$, the adjusted count is 1 of 10 subsamples in the wind tunnel (Wilson $[0.018, 0.404]$, 0.1 false edges per subsample) and 0 of 10 in the light tunnel ($[0, 0.278]$). In the light tunnel every other configuration in the table is at 10 of 10 except the linear certificate at $k = 3$ (4 of 10); in the wind tunnel every other configuration is at 10 of 10 except the nonparametric certificate at $k = 2$ (1 of 10). The nonparametric certificate certifies 10.1 edges per subsample in the light tunnel, of which about 7 are graph edges and 3 are pairs of V, and its recall of the graph edges is only 0.12 and 0.18. The strict count remains 10 of 10 because the pairs of V are certified, which is the intended behaviour for pairs whose dependence the graph does not explain, but it means that against the graph alone CERT-CD still reports false edges. The baselines are scored on the same subsamples; PC, GES

Table 14 Untouched Causal Chamber recordings. For each group: recordings in which a false edge was reported at some sample size, adjusted for the reference set V (strict, graph edges only, in parentheses), false edges per recording at $n = 4000$ (adjusted) and recall of the graph edges. Wind tunnel walks: 15 recordings, 16 nodes, 26 edges, 31 reference pairs on average; pressure control: 1, 11 nodes, 18 edges; light tunnel walks: 3, 13 nodes on average, 26 edges, 14 reference pairs.

Method	wind walks (15)			pressure control (1)			light walks (3)		
	ever	/run	rec.	ever	/run	rec.	ever	/run	rec.
CERT-CD nonparam., $k = 3$	10 (14)	1.1	0.25	1 (1)	1.0	0.12	1 (3)	0.3	0.24
CERT-CD linear, $k = 3$	15 (15)	2.2	0.26	1 (1)	1.0	0.12	0 (3)	0.0	0.17
PC, $\alpha = 0.05$	15 (15)	5.9	0.42	1 (1)	2.0	0.25	2 (3)	4.3	0.37
PC, $\alpha = 0.01$	15 (15)	4.9	0.38	1 (1)	1.0	0.25	2 (3)	3.0	0.36
GES-BIC	15 (15)	25.8	0.71	1 (1)	3.0	0.62	3 (3)	2.0	0.78
DirectLiNGAM	15 (15)	24.0	0.62	1 (1)	4.0	0.56	2 (3)	3.0	0.74
PC, nonparam. test, $\alpha = 0.05$	15 (15)	6.0	0.48	1 (1)	5.0	0.31	2 (3)	1.7	0.49
PC, nonparam. test, $\alpha = 0.01$	15 (15)	4.5	0.41	1 (1)	3.0	0.25	2 (3)	2.7	0.48

and DirectLiNGAM were run with their usual settings and were not given a larger conditioning order or a nonparametric test.

Untouched recordings.

Because the choices above were made after seeing the first two recordings, we froze the algorithms and ran them on recordings that had not been used (a pre-specified rule: all 15 random-walk runs of the wind tunnel with at least 8000 rows, the pressure-control experiment of the wind tunnel, and the three light-tunnel walk experiments, 19 recordings; the sizes of the graphs are 10 to 20 nodes and 18 to 39 edges). Each recording was shuffled and split in halves; the algorithms monitored 4000 rows of one half and the reference set V was built on the other half, so that V does not depend on the monitored rows. Table 14 and the text report what happened.

The guarantee did not hold on these recordings either. The nonparametric certificate at $k = 3$ reported a false edge, adjusted for V, in 12 of the 19 recordings (and a strict false edge in 18 of 19), against 18 or 19 of 19 for the comparators. It had the fewest false edges per recording in the wind-tunnel walks (1.1 against 2.2 to 25.8) and, with the linear variant and PC at $\alpha = 0.01$, among the fewest in the other two groups (1.0 and 0.3), but its recall of the graph edges is 0.25, 0.12 and 0.24, below that of every comparator except the linear variant. Counted by precision with V taken as correct, it certified about 92% of its 14.5 edges per recording correctly in the wind-tunnel walks, against 80% for PC at $\alpha = 0.01$, 84% for PC with the same test and about 60% for GES and DirectLiNGAM. We conclude that, on real devices, the certificate behaves as a conservative edge detector with the fewest false edges among the methods compared, not as a procedure with a guaranteed false-edge probability.

Direction.

Among the edges certified by CERT-CD that are true edges, the bivariate agreement certificate decided 30 of 31 in the wind tunnel and 27 of 70 in the light tunnel (the constant-scale class alone 50 of 70), and every decision was correct (30/30 and 27/27). Separately, we applied the certificate to every true directed edge among the varying

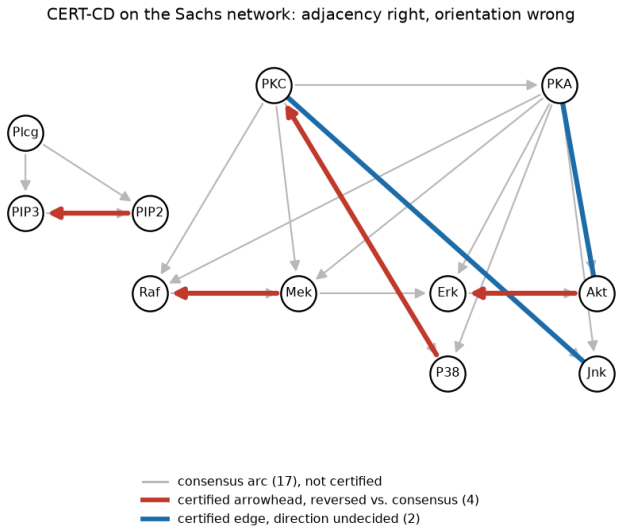


Fig. 4 Sachs data: certified edges and arrows against the consensus network.

variables, each on two random subsamples of 2000 rows (16 and 78 pair-runs for the two devices). The agreement certificate committed on 8 of 16 and 23 of 78 and was right every time (31/31); the constant-scale class alone committed on 8 and 40, and the location-scale class on 9 and 25, also always correctly. The forced accuracy of the two classes is 0.69 and 0.69 in the wind tunnel and 0.74 and 0.81 in the light tunnel, against 0.88 and 1.00 for RECI, 1.00 and 0.92 for IGCI, 0.38 and 0.81 for ANM-HSIC and 0.81 and 0.77 for the pairwise LiNGAM rule. As on the Tübingen pairs, the forced accuracy of the forecasters is not better than a simple rule, and the certificate is valuable for what it declines to decide. The causes in these experiments are uniformly distributed actuators, an easy case for identifiability; the 31 decisions are 47 edges analysed twice, and the 30 and 27 decisions above pool the ten subsamples, so neither set is independent.

4.11 Scope: where the certificates fail

Assumption violations.

We simulated 200 datasets of $n = 3000$ for each of four violations at $\alpha = 0.05$. In-model there were no false edges and no wrong arrows. Violating the separator-size premise gave a false edge in 200/200 runs and false arrows in 26; a latent common cause gave a certified edge in 200/200 runs, which is a MAG adjacency and not an error in that sense, with 38 wrong arrows; nonlinear mechanisms gave false edges in 200/200 runs and 142 wrong arrows; Poisson counts gave false edges in 200/200 runs and 28 wrong arrows. The adjacency certificate is valid only for the linear model with independent errors, and the orientation certificate inside the multivariate algorithm fails badly under nonlinearity.

Sachs protein-signalling data.

The data (Sachs, Perez, Pe'er, Lauffenburger, & Nolan, 2005) are discrete-like, nonlinear and heavy-tailed and sit outside the method's scope, which is why we report them here. On 853 cells (753 after a 100-cell warm-up; Figure 4) CERT-CD certified 6 edges among 55 pairs, all in the consensus skeleton, left 49 pairs undecided, and certified 4 arrowheads, all reversed relative to the consensus network. PC found 7 edges, all in the consensus skeleton. DirectLiNGAM's estimated causal order contradicts the consensus direction on all six pairs CERT-CD certified (recomputed from the stored order). Consensus directions are themselves contested (Mooij, Magliacane, & Claassen, 2020), but the certified arrows are confidently unreliable, and nothing in the output says so. We do not compare with the iterative protocol of Uehara (2026), whose queries, ground truth and interaction model differ from ours.

5 Discussion

What is established.

Under monitoring, the adjacency certificate kept the probability of ever reporting a false edge near 3–4% in random eight-node graphs, where the comparators reached 27% to 100%, and its stated bound is exact in the Gaussian regression model. The orientation certificate commits rarely and, on the Tübingen pairs under a frozen protocol, was right on 11 of the 12 pairs on which it committed.

Limitations.

- *The adjacency bound is asymptotic off the Gaussian model.* The null that a partial regression coefficient is zero follows from d-separation in a linear model with independent errors, but the primitive is exact only for Gaussian regression errors. Section 4.6 measures how far from nominal it is under skewed, heavy-tailed and heteroscedastic errors; we have no finite-sample correction. Conditional independence testing without a model restriction is impossible (Shah & Peters, 2020), so none is to be expected without further assumptions.
- *Power at small samples.* Skeleton recall is 0.21 at $n = 100$ and 0.65 at $n = 200$, and at a single large sample size DirectLiNGAM is more accurate on the same data. The case for CERT-CD is monitoring, not single-sample accuracy.
- *The order k.* Validity needs k at least the size of a separator for every non-adjacent pair; $k = 0$ and $k = 1$ produced false edges in most graphs (Section 4.4), and nothing in the output warns of it.
- *Real physical devices.* On the first two Causal Chamber recordings every method, including CERT-CD with the linear primitive at $k = 2$, reported edges that the graph does not contain in every subsample. The choices that improved the result (the nonparametric test and $k = 3$) were made after seeing these data. On 19 untouched recordings, with the algorithms frozen, CERT-CD still reported a false edge adjusted for the reference set in 12 recordings (a strict false edge in 18), although fewer false edges per recording than the comparators and with recall 0.12 to 0.25 (Section 4.10). The false-edge guarantee should not be read as a property of any real dataset, and it needs k at least the size of the largest separator, which is not known in practice.

- *Orientation at matched coverage.* Ranked by its own confidence, RECI is as accurate as the certificate on the pairs where the certificate commits (Table 12). The certificate's distinct feature is the data-driven stopping rule with an asymptotic statement about better-fit direction, not a higher accuracy.
- *Model-class errors in orientation.* On a family with a bimodal cause and Gaussian noise, the agreement certificate was wrong on all 5 pairs it decided, because a generalised-Gaussian marginal cannot represent the cause.
- *Most pairs stay undecided.* About three quarters of the pairs in the eight-node experiments are undecided at $n = 6400$ because the algorithm never certifies absence, and the orientation certificate decided 15% of the scorable Tübingen pairs. Certificate-only output trades coverage for validity, and a user who needs a complete graph needs a different method.
- *Predictive fit is not causation.* The orientation certificate certifies a better expected log score. It coincides with the causal direction under additive-noise or linear non-Gaussian identifiability and under causal sufficiency; under confounding it need not, and the multivariate algorithm assumes causal sufficiency.
- *The orientation layer inside the multivariate algorithm is the old linear certificate.* Its guarantee is conditional on a linear generalised-Gaussian model and fails outside it (Section 4.11). Conditioning the relative-fit forecasters on the frozen block would remove this, and is not done here.
- *Assumed, not checked.* The linear structural model with independent errors is an assumption. Residual diagnostics on the warm-up prefix can flag gross departures, but they are not part of the guarantee.
- *Small evidence base for orientation.* The Tübingen result rests on 12 committed pairs out of 80 scorable, and 22 of the 102 pairs were too short to score; the confidence interval on the accuracy is wide, and the binomial test treats pairs as independent although some share a dataset. The benchmark was known to us when the certificate was designed (Section 4.9).
- *Domain.* The experiments use continuous data from mostly low-dimensional systems. The Sachs data, which are not of this kind, are reported as out-of-scope behaviour (Section 4.11).

6 Conclusion

Asserting a feature of a causal graph only by rejecting a null that is true when the feature is absent removes the dependence of the false-edge guarantee on faithfulness and makes the guarantee survive monitoring. In simulation within the models, the probability of ever reporting a false edge stays at a few percent where deletion-based methods, including PC with the same test, reach a third to most of the graphs, at the cost of undecided pairs and low recall at small samples. A predictive-fit orientation certificate extends the same discipline to direction; it commits rarely and mostly correctly, but at the same coverage a simple regression-error rule is as accurate, so its contribution is the stopping rule and the stated level. On real physical devices the false-edge guarantee did not hold, though the certificate reported the fewest false

edges of the methods compared. Both certificates are valid only within their models, and the paper reports where they fail.

Competing interests. The author declares no competing interests.

Funding. This research received no specific grant from any funding agency in the public, commercial, or not-for-profit sectors.

Ethics approval. Not applicable. The study involves no human participants or animals. The protein-signalling data analysed in Section 4.11 are publicly available and were published by Sachs et al. (2005).

Data and code availability. All simulation code, experiment scripts, random seeds and figure-generating scripts are available at https://github.com/vinhqdang/SPRINT-CD. The protein-signalling data of Sachs et al. (2005) are redistributed with the `bnlearn` package (Scutari, 2010) and are downloaded by the experiment script rather than included in the repository. The Tübingen cause-effect pairs of Section 4.9 are downloaded on first use from https://webdav.tuebingen.mpg.de/cause-effect/ and are not redistributed.

Author contributions. The author is solely responsible for all aspects of this work.

Use of generative AI. A large language model was used in support of writing this manuscript. The author verified and modified the resulting content and takes full responsibility for the manuscript in its entirety, including its claims, derivations, experiments and references.

Appendix A The safe linear primitive

Setup.

Let $y \in \mathbb{R}^n$ stack the observations of the response X_i, x those of the tested covariate X_j, and $W \in \mathbb{R}^{n\times q}$ the conditioning block with an intercept, $q = |S| + 1$. The model is $y = W\alpha + \beta x + \varepsilon$, $\varepsilon \sim N(0, \sigma^2 I_n)$, with null $\beta = 0$. Let $M = I - W(W^\top W)^{-1}W^\top$ and $\tilde{y} = My$, $\tilde{x} = Mx$, and write $S_{ij} = \tilde{x}^\top \tilde{y}$, $S_{ii} = \tilde{y}^\top \tilde{y}$, $S_{jj} = \tilde{x}^\top \tilde{x}$, so that $\hat{r}^2 = S_{ij}^2/(S_{ii}S_{jj})$.

Marginal likelihoods.

The null is invariant under $y \mapsto \lambda y + Wc$, $\lambda > 0$, and the right-Haar prior is $\pi(\alpha, \sigma) \propto \sigma^{-1}$. Integrating α and σ gives $m_0(y) = C_{n,q} S_{ii}^{-(n-q)/2}$. Under the alternative, put $\beta \mid \sigma \sim N(0, g\sigma^2)$. Completing the square in β,

$$m_1(y) = C_{n,q}(1 + gS_{jj})^{-1/2}\Big(S_{ii} - \frac{S_{ij}^2}{S_{jj} + g^{-1}}\Big)^{-(n-q)/2},$$

and with $G = gS_{jj}$, $\kappa = G/(1+G)$,

$$E_n = \frac{m_1}{m_0} = (1+G)^{-1/2}\big(1 - \kappa \hat{r}_n^2\big)^{-(n-q)/2},$$

which is (3). The constant g is fixed before the monitored data arrive, from a discarded warm-up prefix.

Proof of Proposition 4.

A Bayes factor with a right-Haar prior on a group acting on the nuisance parameters of both hypotheses equals the Bayes factor of the maximal invariant (Berger, Pericchi, & Varshavsky, 1998; Hendriksen et al., 2021). The maximal invariant's law under the null does not depend on (α, σ), so E_n is a ratio of a density to its null density in the filtration $(\mathcal{G}_n)$ generated by the maximal invariants, hence a test martingale there. Ville's inequality gives $\mathbb{P}(\exists n : E_n \geq 1/\alpha) \leq \alpha$ under every Gaussian null.

The raw-observation filtration.

(E_n) is not a martingale in the filtration of the raw observations. Take $S = \emptyset$, $g = 0.25$ and $X_{1:2} = Y_{1:2} = (-1, 1)$, with the Gram matrix centred. Then $E_2 = 1$, while under an i.i.d. standard Gaussian independence null the exact conditional expectation is $\mathbb{E}[E_3 \mid \mathcal{F}_2] \approx 1.083 > 1$. The guarantee is therefore the Ville bound on the crossing event, which is what Propositions 3 and Theorem 7 use; it is not a claim about $\mathbb{E}[E_\tau]$ at arbitrary stopping times of the raw filtration.

Appendix B The linear orientation certificate inside the multivariate algorithm

Inside the streaming algorithm, direction is certified after the pair's adjacency, with a sequential universal-inference ratio (Wasserman, Ramdas, & Balakrishnan, 2020). For a pair (i, j) certified at the stopping time T, let Z_{ij} be a block of conditioning variables fixed at T. The numerator is a generalised-Gaussian linear factorisation $x \to y$ given Z_{ij} fitted sequentially on the past after T, with shape in $[0.9, 8]$ and a continuation across the Gaussian boundary of the shape grid; the denominator is the supremum over the reversed linear factorisation. The observations before T are discarded. Conditionally on $\mathcal{F}_T$ the later observations are i.i.d. and Z_{ij} is a fixed set, so the usual universal-inference argument gives an e-process for the null that the reversed linear factorisation holds; the union bound over the $d(d-1)$ ordered pairs uses the fixed family and does not depend on which hypotheses the run examines. The guarantee is conditional on the linear generalised-Gaussian model describing the pair given Z_{ij}, and it fails outside it (Section 4.11). A direction is reported after at least 80 post-certification observations.

Appendix C Inverting the primitive into a confidence sequence

Replacing y by $y - \beta_0 x$ gives the process for the null $\beta = \beta_0$, and the set of β_0 not rejected at $1/\alpha$ is the interval

$$\hat{\beta} \pm \frac{1}{S_{jj}} \sqrt{\frac{\tau(S_{ii}S_{jj} - S_{ij}^2)}{1-\tau}}, \qquad \tau = \frac{1}{\kappa}\Big[1 - \Big(\frac{\alpha}{\sqrt{1+G}}\Big)^{2/(n-q)}\Big],$$

valid for $0 < \tau < 1$. Since $\alpha < 1$, the term in brackets is positive and $\tau > 0$ for every $n > q$; when $\tau \geq 1$ the set is the whole line. The closed form is checked in the test suite against a numerical inversion over a grid of β_0.

Appendix D Additional experimental detail

Hyperparameters.

Linear primitive: warm-up of 50 discarded rows; prior standard deviation 0.5 on the standardised scale ($g = 0.25$); $k = 2$; adjacency share $\alpha_A/\alpha = 0.5$; batches of 100 rows in all monitoring experiments except the faithfulness experiment, which uses 250. Nonparametric primitive: warm-up of 100 discarded rows, $k = 2$ ($k = 3$ on the Causal Chambers), per-pair level $\alpha_A/\binom{d}{2}$ with $\alpha_A = 0.05$, ridge penalty 1, refits when the sample has grown by a factor 1.25, batches of 100 rows. Pair orientation: $\alpha = 0.05$, ρ tuned to $t^* = 200$, a 60-observation minimum, warm-up $\max(100, \lceil 0.2n \rceil)$, refits every 50 scored rows.

Calibration of the linear primitive under six error laws.

Table D1 reports the realised probability of ever crossing for the linear primitive with homoscedastic errors, 2000 replications, monitored to $n = 800$ after every tenth observation, prior scale 0.4 ($g = 0.16$). Every entry is at or below its nominal level; the largest realised-to-nominal ratio is 0.95 (shifted exponential, $\alpha = 0.01$) and the median is 0.48.

Table D1 Realised probability of ever crossing the threshold, linear primitive, six error laws.

error law	$\alpha = 0.20$	$\alpha = 0.10$	$\alpha = 0.05$	$\alpha = 0.01$
Gaussian	0.091	0.046	0.023	0.003
Laplace	0.092	0.047	0.024	0.003
uniform	0.107	0.052	0.031	0.006
Student-t_5	0.093	0.044	0.020	0.004
shifted exponential	0.101	0.051	0.028	0.010
two-component mixture	0.097	0.049	0.028	0.006

Comparison rules.

IGCI is the entropy form with a uniform reference (infer $x \to y$ if the spacing entropy of y is smaller than that of x after min–max scaling) (Daniušis et al., 2010). We re-implemented the published rules in a few lines each and did not tune them. RECI (Blöbaum, Janzing, Washio, Shimizu, & Schölkopf, 2018) compares the mean squared error of a cubic polynomial regression in each direction after min–max scaling. ANM-HSIC fits a cubic spline regression in each direction and prefers the direction in which the residual has the smaller HSIC with the regressor (Gaussian kernels, median bandwidth, at most 500 rows). The pairwise LiNGAM rule is the first-order likelihood-ratio statistic of Hyvärinen and Smith (2013) with the sign set by the sample kurtosis.

Reproduction.

The experiment scripts are in `experiments/`, named by experiment number, and the figure and summary scripts in `scripts/`; seeds are fixed in each script and results are written to `results/` as JSON. Scripts for the monitoring, Tübingen, pair-orientation, scaling, nonlinear and Causal Chambers experiments were committed before they were run.

References

Berger, J.O., Pericchi, L.R., Varshavsky, J.A. (1998). Bayes factors and marginal distributions in invariant situations. *Sankhyā: The Indian Journal of Statistics Series A*, *60*(3), 307–321,

Blöbaum, P., Janzing, D., Washio, T., Shimizu, S., Schölkopf, B. (2018). Cause-effect inference by comparing regression errors. *Proceedings of the 21st international conference on artificial intelligence and statistics.*

Chernozhukov, V., Chetverikov, D., Demirer, M., Duflo, E., Hansen, C., Newey, W., Robins, J. (2018). Double/debiased machine learning for treatment and structural parameters. *The Econometrics Journal*, *21*(1), C1–C68,

Choe, Y.J., & Ramdas, A. (2024). Comparing sequential forecasters. *Operations Research*, *72*(4), ,

Colombo, D., Maathuis, M.H., Kalisch, M., Richardson, T.S. (2012). Learning high-dimensional directed acyclic graphs with latent and selection variables. *The Annals of Statistics*, *40*(1), 294–321, https://doi.org/10.1214/11-AOS940

Csillag, D., Struchiner, C.J., Goedert, G.T. (2025). Prediction-powered e-values. *Proceedings of the 42nd international conference on machine learning.* (arXiv:2502.04294)

Daniušis, P., Janzing, D., Mooij, J., Zscheischler, J., Steudel, B., Zhang, K., Schölkopf, B. (2010). Inferring deterministic causal relations. *Proceedings of the 26th conference on uncertainty in artificial intelligence.*

Darling, D.A., & Robbins, H. (1967). Confidence sequences for mean, variance, and median. *Proceedings of the National Academy of Sciences*, *58*(1), 66–68,

Ding, Z., & Zhang, X.-P. (2025). *GaussDetect-LiNGAM: Causal direction identification without Gaussianity test.* (Preprint, arXiv:2512.03428)

Gamella, J.L., Peters, J., Bühlmann, P. (2025). Causal chambers as a real-world physical testbed for AI methodology. *Nature Machine Intelligence*, , https://doi.org/10.1038/s42256-024-00964-x

Grünwald, P., de Heide, R., Koolen, W.M. (2024). Safe testing. *Journal of the Royal Statistical Society Series B: Statistical Methodology*, *86*(5), 1091–1128, https://doi.org/10.1093/jrsssb/qkae011

Grünwald, P., Henzi, A., Lardy, T. (2024). Anytime-valid tests of conditional independence under model-X. *Journal of the American Statistical Association*, *119*(546), ,

Hendriksen, A., de Heide, R., Grünwald, P. (2021). Optional stopping with Bayes factors: A categorization and extension of folklore results, with an application to invariant situations. *Bayesian Analysis*, *16*(3), 961–989, https://doi.org/10.1214/20-BA1234

Howard, S.R., Ramdas, A., McAuliffe, J., Sekhon, J. (2021). Time-uniform, nonparametric, nonasymptotic confidence sequences. *The Annals of Statistics*, *49*(2), 1055–1080, https://doi.org/10.1214/20-AOS1991

Hoyer, P.O., Janzing, D., Mooij, J.M., Peters, J., Schölkopf, B. (2009). Nonlinear causal discovery with additive noise models. *Advances in neural information processing systems 21* (pp. 689–696).

Hyvärinen, A., & Smith, S.M. (2013). Pairwise likelihood ratios for estimation of non-Gaussian structural equation models. *Journal of Machine Learning Research*, *14*, 111–152,

Immer, A., Schultheiss, C., Vogt, J.E., Schölkopf, B., Bühlmann, P., Marx, A. (2023). On the identifiability and estimation of causal location-scale noise models. *Proceedings of the 40th international conference on machine learning.*

Li, J., & Wang, Z.J. (2009). Controlling the false discovery rate of the association/causality structure learned with the PC algorithm. *Journal of Machine Learning Research*, *10*, 475–514,

Lindon, M., Ham, D.W., Tingley, M., Bojinov, I. (2026). Anytime-valid inference in linear models with applications to regression-adjusted causal inference. *Journal of the American Statistical Association*, 1–12, https://doi.org/10.1080/01621459.2026.2692052

Lundborg, A.R., Shah, R.D., Peters, J. (2022). Conditional independence testing in Hilbert spaces with applications to functional data analysis. *Journal of the Royal Statistical Society Series B*, *84*(5), 1821–1850,

Marx, A., & Vreeken, J. (2017). Telling cause from effect using MDL-based local and global regression. *2017 ieee international conference on data mining.*

Mooij, J.M., Magliacane, S., Claassen, T. (2020). Joint causal inference from multiple contexts. *Journal of Machine Learning Research*, *21*(99), 1–108,

Mooij, J.M., Peters, J., Janzing, D., Zscheischler, J., Schölkopf, B. (2016). Distinguishing cause from effect using observational data: Methods and benchmarks. *Journal of Machine Learning Research*, *17*(32), 1–102,

Peters, J., & Bühlmann, P. (2014). Identifiability of Gaussian structural equation models with equal error variances. *Biometrika*, *101*(1), 219–228, https://doi.org/10.1093/biomet/ast043

Podkopaev, A., & Ramdas, A. (2023). Sequential predictive two-sample and independence testing. *Advances in neural information processing systems* (Vol. 36, pp. 53275–53307).

Prakash, S., Xia, F., Erosheva, E. (2024). *A diagnostic tool for functional causal discovery.* (Preprint, arXiv:2406.07787)

Prakash, S., Xia, F., Erosheva, E.A. (2026). *Causal discovery via statistical power.* (Preprint, arXiv:2605.13550)

Ramdas, A., Grünwald, P., Vovk, V., Shafer, G. (2023). Game-theoretic statistics and safe anytime-valid inference. *Statistical Science*, *38*(4), 576–601, https://doi.org/10.1214/23-STS894

Ramsey, J., Spirtes, P., Zhang, J. (2006). Adjacency-faithfulness and conservative causal inference. *Proceedings of the 22nd conference on uncertainty in artificial intelligence* (pp. 401–408).

Reisach, A.G., Seiler, C., Weichwald, S. (2021). Beware of the simulated DAG! causal discovery benchmarks may be easy to game. *Advances in neural information processing systems* (Vol. 34).

Robbins, H. (1970). Statistical methods related to the law of the iterated logarithm. *The Annals of Mathematical Statistics*, *41*(5), 1397–1409,

Robins, J.M., Scheines, R., Spirtes, P., Wasserman, L. (2003). Uniform consistency in causal inference. *Biometrika*, *90*(3), 491–515, https://doi.org/10.1093/biomet/90.3.491

Sachs, K., Perez, O., Pe'er, D., Lauffenburger, D.A., Nolan, G.P. (2005). Causal protein-signaling networks derived from multiparameter single-cell data. *Science*, *308*(5721), 523–529, https://doi.org/10.1126/science.1105809

Schkoda, D., & Drton, M. (2025). Goodness-of-fit tests for linear non-Gaussian structural equation models. *Biometrika*, *112*(4), asaf046,

Scutari, M. (2010). Learning Bayesian networks with the bnlearn R package. *Journal of Statistical Software*, *35*(3), 1–22, https://doi.org/10.18637/jss.v035.i03

Shaer, S., Maman, G., Romano, Y. (2023). Model-X sequential testing for conditional independence via testing by betting. *Proceedings of the 26th international conference on artificial intelligence and statistics.*

Shah, R.D., & Peters, J. (2020). The hardness of conditional independence testing and the generalised covariance measure. *The Annals of Statistics*, *48*(3), 1514–1538, https://doi.org/10.1214/19-AOS1857

Shaska, J., & Mitra, U. (2025). *Causal link discovery with unequal edge error tolerance.* (Preprint, arXiv:2507.21570)

Shimizu, S., Hoyer, P.O., Hyvärinen, A., Kerminen, A. (2006). A linear non-Gaussian acyclic model for causal discovery. *Journal of Machine Learning Research*, *7*, 2003–2030,

Shimizu, S., Inazumi, T., Sogawa, Y., Hyvärinen, A., Kawahara, Y., Washio, T., ... Bollen, K. (2011). DirectLiNGAM: A direct method for learning a linear non-Gaussian structural equation model. *Journal of Machine Learning Research*, *12*, 1225–1248,

Shimizu, S., & Kano, Y. (2008). Use of non-normality in structural equation modeling: Application to direction of causation. *Journal of Statistical Planning and Inference*, *138*(11), 3483–3491, https://doi.org/10.1016/j.jspi.2006.01.017

Spirtes, P., Glymour, C., Scheines, R. (2000). *Causation, prediction, and search* (2nd ed.). Cambridge, MA: MIT Press.

Spirtes, P., Meek, C., Richardson, T. (1995). Causal inference in the presence of latent variables and selection bias. *Proceedings of the 11th conference on uncertainty in artificial intelligence* (pp. 499–506).

Strobl, E.V., Spirtes, P.L., Visweswaran, S. (2019). Estimating and controlling the false discovery rate of the PC algorithm using edge-specific p-values. *ACM Transactions on Intelligent Systems and Technology*, *10*(5), 1–37, https://doi.org/10.1145/3351342

Tagasovska, N., Vatter, T., Chavez-Demoulin, V. (2020). Distinguishing cause from effect using quantiles: bivariate quantile causal discovery. *Proceedings of the 37th international conference on machine learning.*

Uehara, E. (2026). *Iterative causal discovery: Per-edge impossibility certificates, tier-aware oracle queries, and the* $1+k$ *lower bound.* (Preprint, arXiv:2605.27477)

Uhler, C., Raskutti, G., Bühlmann, P., Yu, B. (2013). Geometry of the faithfulness assumption in causal inference. *The Annals of Statistics*, *41*(2), 436–463, https://doi.org/10.1214/12-AOS1080

Ville, J. (1939). *Étude critique de la notion de collectif.* Paris: Gauthier-Villars.

Vovk, V., & Wang, R. (2021). E-values: Calibration, combination and applications. *The Annals of Statistics*, *49*(3), 1736–1754, https://doi.org/10.1214/20-AOS2020

Vuong, Q.H. (1989). Likelihood ratio tests for model selection and non-nested hypotheses. *Econometrica*, *57*(2), 307–333,

Wang, R., & Ramdas, A. (2022). False discovery rate control with e-values. *Journal of the Royal Statistical Society Series B: Statistical Methodology*, *84*(3), 822–852, https://doi.org/10.1111/rssb.12489

Wasserman, L., Ramdas, A., Balakrishnan, S. (2020). Universal inference. *Proceedings of the National Academy of Sciences*, *117*(29), 16880–16890, https://doi.org/10.1073/pnas.1922664117

Waudby-Smith, I., Arbour, D., Sinha, R., Kennedy, E.H., Ramdas, A. (2024). Time-uniform central limit theory and asymptotic confidence sequences. *The Annals of Statistics*, *52*(6), 2613–2640,

Waudby-Smith, I., Kennedy, E.H., Ramdas, A. (2023). Distribution-uniform anytime-valid sequential inference and the Robbins–Siegmund distributions. *arXiv preprint arXiv:2311.03343*, ,

Xu, S., Mian, O.A., Marx, A., Vreeken, J. (2022). Inferring cause and effect in the presence of heteroscedastic noise. *Proceedings of the 39th international conference on machine learning.*

Zhang, J., & Spirtes, P. (2008). Detection of unfaithfulness and robust causal inference. *Minds and Machines*, *18*(2), 239–271, https://doi.org/10.1007/s11023-008-9096-4